

```
In [ ]: # Cell only for Colab

!pip install kagglehub

import kagglehub
import os
import shutil

blood_cell_cancer_all_4class_path = kagglehub.dataset_download('mohammadimir/blood-cell-cancer-all-4class')

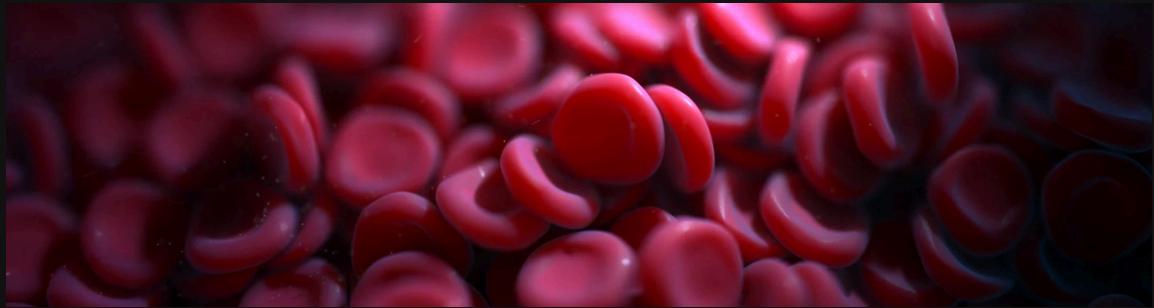
for dir in ["kaggle", "kaggle/input", "kaggle/working", "kaggle/working/model_weights"]:
    if not os.path.exists(dir):
        os.mkdir(dir)

target_dir = 'kaggle/input'

file_names = os.listdir(blood_cell_cancer_all_4class_path)

for file_name in file_names:
    shutil.move(os.path.join(blood_cell_cancer_all_4class_path, file_name), target_dir)
```

血腥 🩸 血液细胞癌预测 - 分类与分割 by Alexander Daniel Rios



Overview

急性淋巴细胞白血病(ALL)是一种影响血液和骨髓的癌症(即在骨骼中产生血液细胞的海绵状组织)。

急性淋巴细胞白血病是儿童中最常见的癌症，治疗效果良好。急性淋巴细胞白血病也可发生在成年人身上，但治愈的机会大大降低。

Causes

急性淋巴细胞白血病发生于骨髓细胞发生突变时。细胞的DNA包含指导细胞生长的指令。

do. Normally, the DNA tells the cell to grow at a set rate and to die at a set time. In acute lymphocytic leukemia, the mutations tell the bone marrow cell to continue growing and dividing.

When this happens, blood cell production becomes out of control. The bone marrow produces immature cells that develop into leukemic white blood cells called lymphoblasts. These abnormal cells are unable to function properly, and they can build up and crowd out healthy cells.

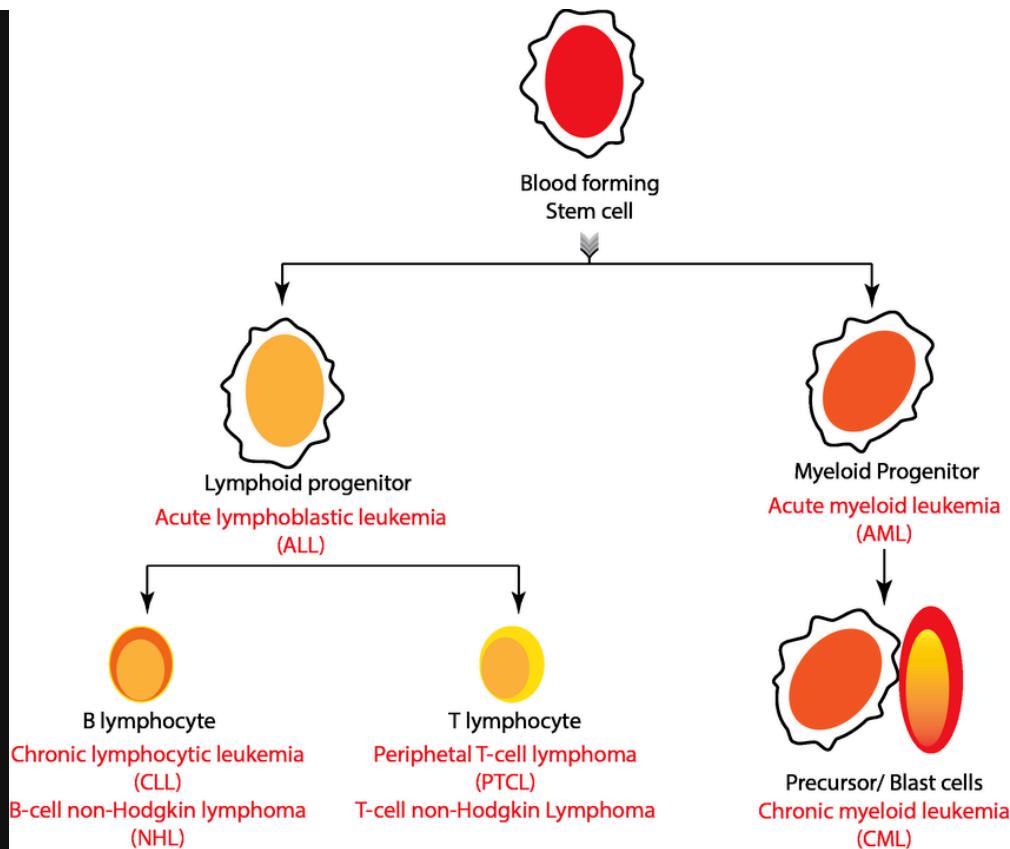
ALL Subtypes

Doctors classify acute lymphoblastic leukemia (ALL) into subtypes by using various tests. It's important to get an accurate diagnosis since your subtype plays a large part in deciding the type of treatment you'll receive. Depending on your ALL subtype, the doctor will determine

- The type of drug combination needed for your treatment
- The length of time you'll need to be in treatment
- Other types of treatment that may be needed to achieve the best outcomes

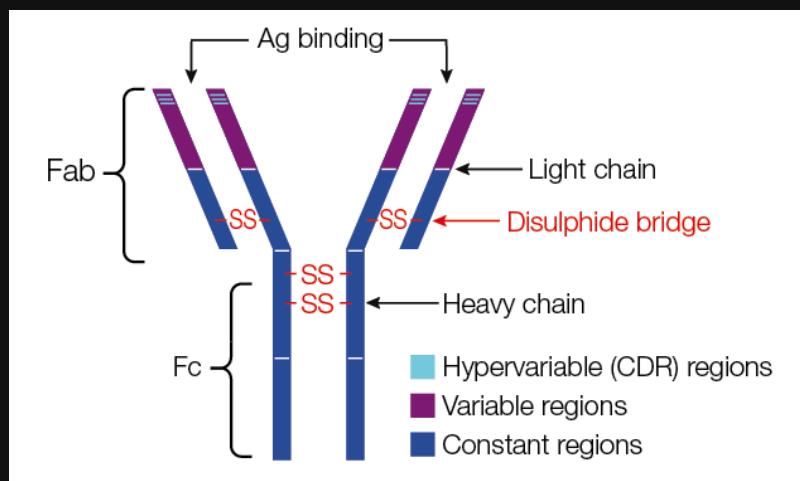
Based on immunophenotyping of the leukemia cell, the World Health Organization (WHO) classifies ALL into two main subtypes.

- **B-cell lymphoblastic leukemia/lymphoma:** This subtype begins in immature cells that would normally develop into B-cell lymphocytes. This is the most common ALL subtype. Among adults, B-cell lineage represents 75 percent of cases.
- **T-cell lymphoblastic leukemia:** This subtype of ALL originates in immature cells that would normally develop into T-cell lymphocytes. This subtype is less common, and it occurs more often in adults than in children. Among adults, T-cell lineage represents about 25 percent of cases.



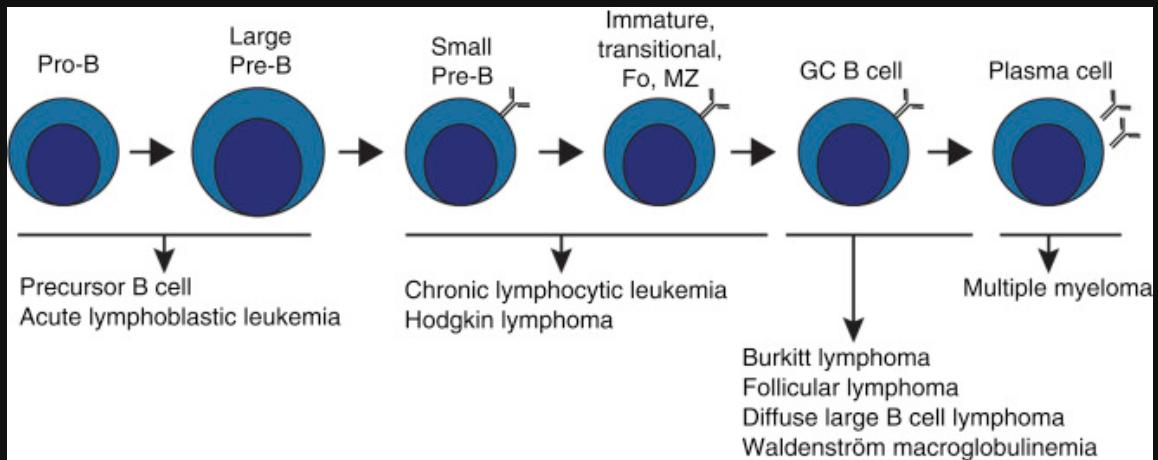
The final task of the lymphocytes (B cells) developed in the bone marrow (BM) is the production of antigen-specific immunoglobulins (IGS), which function as antibodies. IGS are proteins secreted by or present on the surface of B cells, assembled from identical pairs of heavy (H) and light (L) chains.

Antibody structure



Before being capable of producing Ag-specific IGS, B cells must undergo a number of transformations, first in the BM and subsequently in the lymph nodes (LN). The first stages of B-cell development occur in the BM, where pro-B cells first rearrange the Ig H chain gene to become pre-B cells. Pre-B cells continue this somatic recombination process by rearranging the L chain to become immature B cells, expressing IgM on their surface.

B-cell development



About the dataset

In this project I used the following dataset: [Blood cell cancer all 4 class](#)

You can download this dataset with the following code:

```
!pip install kagglehub

import kagglehub
blood_cell_cancer_all_4class_path =
kagglehub.dataset_download('mohammadamireshraghi/blood-cell-
cancer-all-4class')
```

The images of this dataset were prepared in the bone marrow laboratory of Taleqani Hospital (Tehran, Iran).

This dataset consisted of **3242 PBS images** from 89 patients suspected of ALL, whose blood samples were prepared and stained by skilled laboratory staff. This dataset is divided into two classes *benign* and *malignant*. The former comprises hematogenous, and the latter is the ALL group with three subtypes of malignant lymphoblasts: **Early Pre-B**, **Pre-B**, and **Pro-B ALL**.

All the images were taken by using a Zeiss camera in a microscope with a 100x magnification and saved as JPG files. A specialist using the flow cytometry tool made the definitive determination of the types and subtypes of these cells.

Table of Contents

1. [Installing Packages](#)
2. [Importing Packages](#)
3. [Setting for Reproducibility](#)

4. Gathering Image Pathfiles
 5. EDA: Visualization of the Image Dataset
 6. Extracting the Masks from the Images
 7. Creating Training, Validation, and Testing Image Sets
 8. Handling Imbalance in the Image Dataset
 9. Diagnosing Acute Lymphoblastic Leukemia (ALL) by Classifying and Segmenting Images of Blood Cells Affected by Cancer
 10. Image Classification
 11. Image Segmentation
 12. Image Segmentation and Classification
 13. Let's Take a Look at the Predictions
-

Installing packages

```
In [ ]: !pip install mplcyberpunk
```

Importing packages

```
In [ ]: import numpy as np
import pandas as pd
import os
import random

import keras
import tensorflow as tf
from tensorflow.keras import backend as K
from tensorflow.keras.backend import epsilon
from tensorflow.keras.utils import load_img
from tensorflow.data.experimental import cardinality
from tensorflow.image import (random_flip_left_right,
                               random_flip_up_down,
                               random_brightness,
                               random_contrast,
                               random_jpeg_quality,
                               random_saturation,
                               resize)

from tensorflow.keras import Model, Input
from tensorflow.keras.layers import (Conv2D,
                                      Dense,
                                      Dropout,
                                      Flatten,
                                      GlobalAveragePooling2D,
                                      MaxPooling2D,
                                      BatchNormalization,
                                      ReLU,
```

```

        Conv2DTranspose,
        Concatenate)

from tensorflow.keras.regularizers import L2
from tensorflow.keras.losses import CategoricalCrossentropy, Dice
from tensorflow.keras.metrics import AUC, BinaryIoU, BinaryAccuracy
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import Callback, EarlyStopping, ReduceLROnPla
from tensorflow.keras.applications import VGG16
from tensorflow.keras.utils import plot_model

from sklearn.model_selection import train_test_split
from sklearn.utils.class_weight import compute_class_weight
from sklearn.metrics import f1_score, accuracy_score, precision_score, recal

import cv2

from IPython.display import clear_output

import matplotlib.pyplot as plt
import mplcyberpunk
plt.style.use("cyberpunk")

```

Setting for reproducibility

```
In [ ]: seed_value = 42
os.environ['PYTHONHASHSEED'] = str(seed_value)
random.seed(seed_value)
np.random.seed(seed_value)
```

Gathering image pathfiles

```
In [ ]: imgs_path = []
classes = []

data_dir = "/kaggle/input/blood-cell-cancer-all-4class/Blood cell Cancer [AL
for dirpath, dirnames, filenames in os.walk(data_dir):
    if len(classes) == 0:
        classes = dirnames
    for filename in filenames:
        imgs_path.append(os.path.join(dirpath, filename))
```

From images metadata to Pandas DataFrame

```
In [ ]: def get_class(x):
    for class_ in classes:
        if class_ in x:
            return class_
```

```

df_imgs = pd.DataFrame(data=imgs_path, columns=["pathfiles"])
df_imgs["type_cell"] = df_imgs.map(get_class)

widths = []
heights = []
for row in df_imgs.iterrows():
    h, w = load_img(row[1].pathfiles).size
    widths.append(w)
    heights.append(h)

df_imgs["width"] = widths
df_imgs["height"] = heights

df_imgs

```

Out []:

	pathfiles	type_cell	width	height
0	/kaggle/input/blood-cell-cancer-all-4class/Blo...	[Malignant] early Pre-B	768	1024
1	/kaggle/input/blood-cell-cancer-all-4class/Blo...	[Malignant] early Pre-B	768	1024
2	/kaggle/input/blood-cell-cancer-all-4class/Blo...	[Malignant] early Pre-B	768	1024
3	/kaggle/input/blood-cell-cancer-all-4class/Blo...	[Malignant] early Pre-B	768	1024
4	/kaggle/input/blood-cell-cancer-all-4class/Blo...	[Malignant] early Pre-B	768	1024
...
3237	/kaggle/input/blood-cell-cancer-all-4class/Blo...	Benign	768	1024
3238	/kaggle/input/blood-cell-cancer-all-4class/Blo...	Benign	768	1024
3239	/kaggle/input/blood-cell-cancer-all-4class/Blo...	Benign	768	1024
3240	/kaggle/input/blood-cell-cancer-all-4class/Blo...	Benign	1920	2560
3241	/kaggle/input/blood-cell-cancer-all-4class/Blo...	Benign	768	1024

3242 rows × 4 columns

Setting for normalization of image sizes

In []: df_imgs.width.value_counts()

Out []: width
768 3196
1920 46
Name: count, dtype: int64

The common factors between 768 and 1920 are [1, 2, 3, 4, 6, 8, 12, 16, 24, 32, 48, 64, 96, 128, 192, 384], therefore we will resize these images to 192.

In []: df_imgs.height.value_counts()

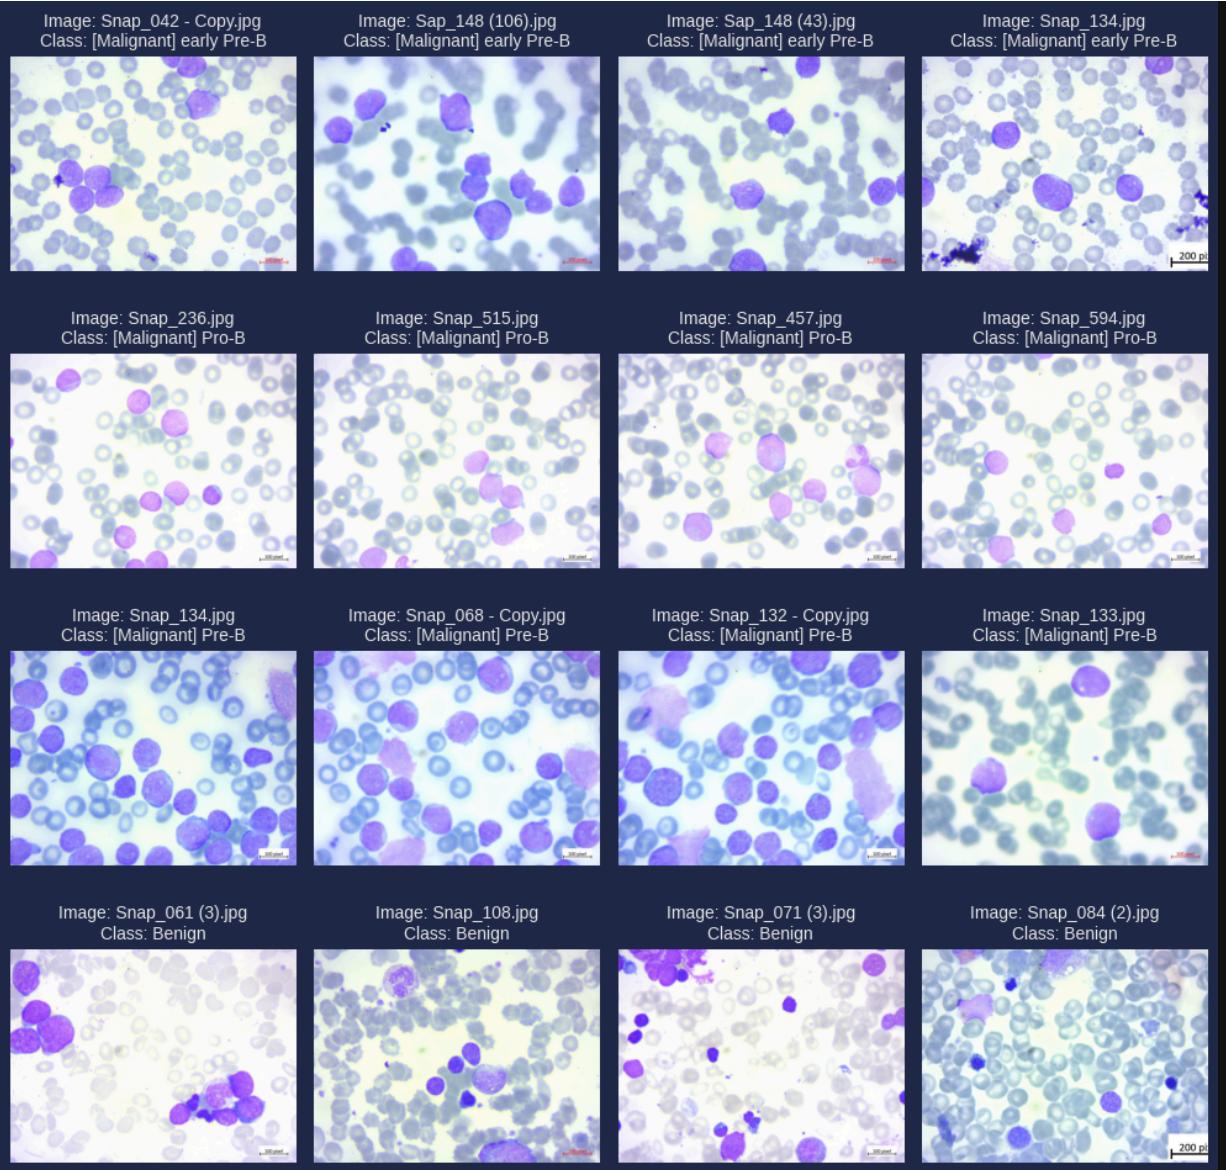
```
Out[ ]: heigth  
       1024      3196  
       2560       46  
Name: count, dtype: int64
```

The common factors between 1024 and 2560 are [1, 2, 4, 8, 16, 32, 64, 128, 256, 512], therefore we will resize these images to 256.

```
In [ ]: img_shape = (192, 256)
```

EDA: Visualization of the image dataset

```
In [ ]: def random_img(df, cell=None, n=4):  
    if cell:  
        aux = df[df["type_cell"] == cell]  
    else:  
        aux = df  
    return aux.loc[random.sample(list(aux.index), n)]  
  
plt.subplots(4, 4, figsize=(10, 10))  
  
for x, class_ in enumerate(classes):  
    imgs = random_img(df_imgs, cell=class_, n=4)  
    for i, row in enumerate(imgs.iterrows()):  
        plt.subplot(4, 4, (x*4)+(i+1))  
        plt.imshow(load_img(row[1].pathfiles))  
        plt.title(f"Image: {row[1].pathfiles.split('/')[-1]}\nClass: {class_}")  
        plt.axis("off")  
plt.tight_layout()
```



Extract the masks from the images

```
In [ ]: def get_image(pathfile):
    return cv2.cvtColor(cv2.imread(pathfile), cv2.COLOR_BGR2RGB)

def RGB2LAB(image):
    image_lab = cv2.cvtColor(image, cv2.COLOR_RGB2LAB)
    l, a, b = cv2.split(image_lab)
    return l, a, b

def get_mask(image):
    l, a, b = RGB2LAB(image)
    a_blur = cv2.GaussianBlur(a, (19, 19), 0)
    _, thresh_img = cv2.threshold(a_blur, 200, 255, cv2.THRESH_BINARY+cv2.THRESH_OTSU)
    kernel = np.ones((1, 1), np.uint8)
    return cv2.morphologyEx(thresh_img, op=cv2.MORPH_CLOSE, kernel=kernel, i
```

```
def apply_mask(image, mask):
    return cv2.bitwise_and(image, image, mask=mask)

plt.subplots(4, 3, figsize=(10, 10))
imgs = random_img(df_imgs, n=4)

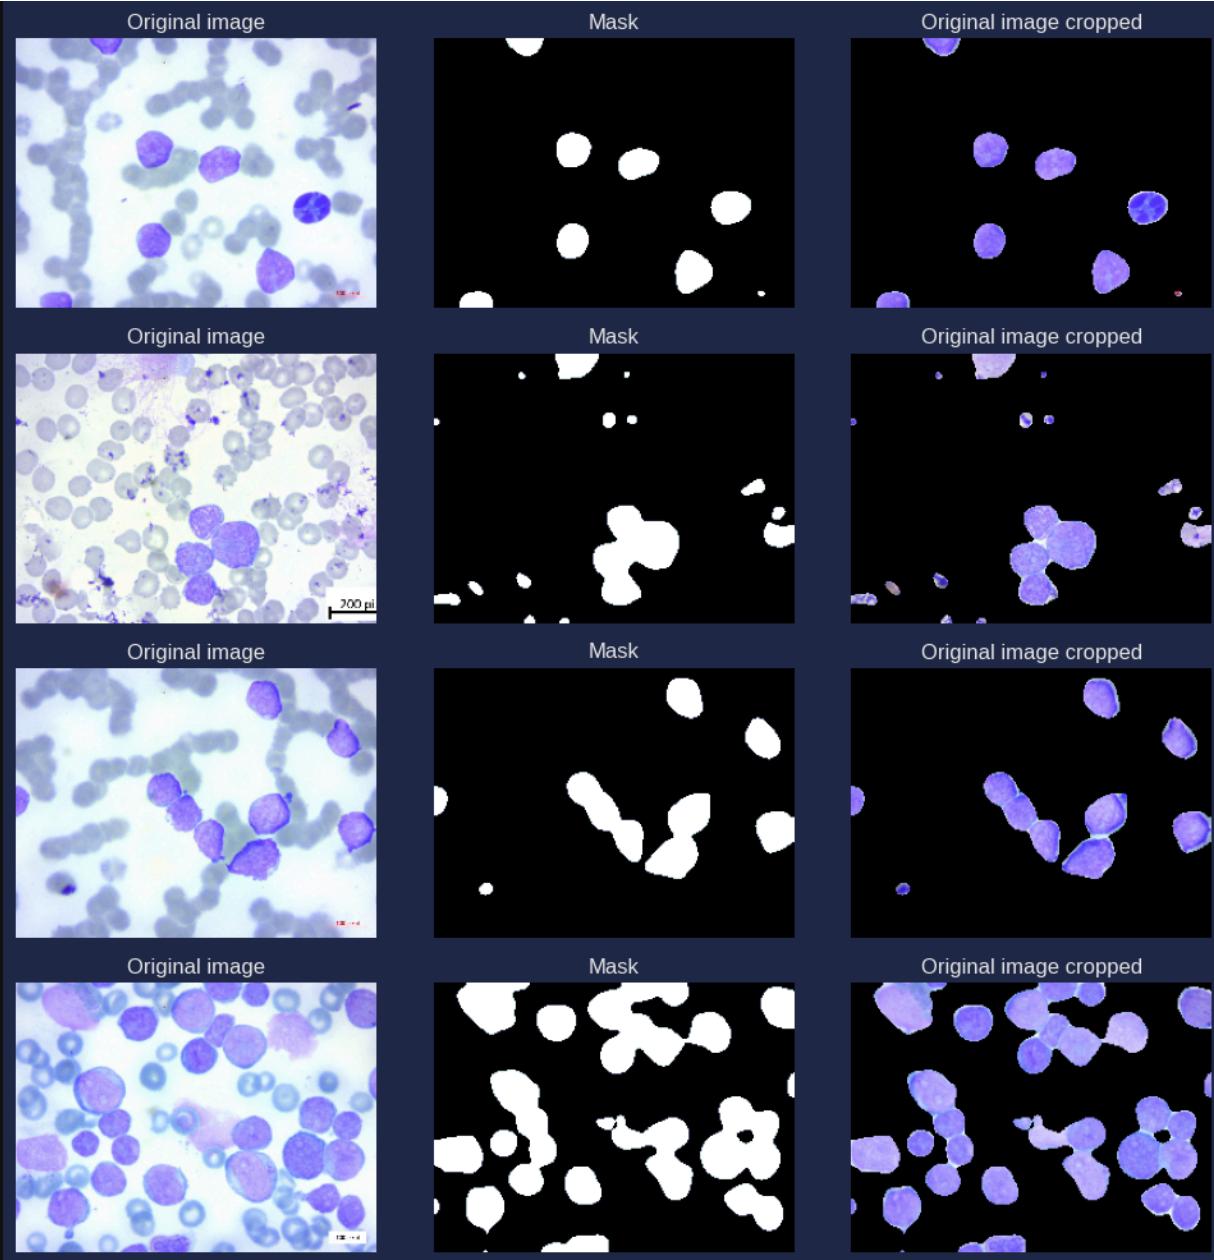
for i, row in enumerate(imgs.iterrows()):
    # Read image
    image = get_image(pathfile=row[1].pathfiles)
    image = resize(image, size=img_shape, method="nearest", antialias=True)
    mask = get_mask(image)
    image_cropped = apply_mask(image, mask)

    plt.subplot(4, 3, (i*3)+1)
    plt.imshow(image)
    plt.axis("off")
    plt.title("Original image")

    plt.subplot(4, 3, (i*3)+2)
    plt.imshow(mask, cmap="bone")
    plt.axis("off")
    plt.title("Mask")

    plt.subplot(4, 3, (i*3)+3)
    plt.imshow(image_cropped)
    plt.axis("off")
    plt.title("Original image cropped");

plt.tight_layout()
plt.show()
```



Creating training, validation and testing image set

```
In [ ]: df_full_train, df_test = train_test_split(df_imgs, test_size=0.15, stratify=df_imgs['type'])
df_train, df_val = train_test_split(df_full_train, test_size=1-(7/8.5), stra
```

Creating TFRecord files for contain our data

```
In [ ]: def createTFRecord(pathfile, dataset, size_img):
    #options = tf.io.TFRecordOptions(compression_type="GZIP", compression_le
    writer = tf.io.TFRecordWriter(pathfile,
```

```

        #options=options
    )
    cnt_written_img = 0
    for index, row in dataset.iterrows():
        image = get_image(row.pathfiles)
        image = resize(image, size=img_shape, method="nearest", antialias=True)
        mask = np.array(get_mask(image))

        class_ = np.argmax(np.array(classes) == row.type_cell)
        image_data, mask_data = image.tobytes(), mask.tobytes()

        example = tf.train.Example(features=tf.train.Features(feature={
            'image': tf.train.Feature(bytes_list=tf.train.BytesList(value=[image_data])),
            'mask': tf.train.Feature(bytes_list=tf.train.BytesList(value=[mask_data])),
            'class': tf.train.Feature(int64_list=tf.train.Int64List(value=[class_]))
        }))

        writer.write(example.SerializeToString())
        cnt_written_img += 1
    return cnt_written_img

def parse(feature):
    features = tf.io.parse_single_example(
        feature,
        features={
            'image': tf.io.FixedLenFeature([], tf.string),
            'mask': tf.io.FixedLenFeature([], tf.string),
            'class': tf.io.VarLenFeature(tf.int64),
        })
    image = tf.reshape(tf.io.decode_raw(features['image'], out_type=tf.uint8))
    mask = tf.reshape(tf.io.decode_raw(features['mask'], out_type=tf.uint8))
    class_ = tf.sparse.to_dense(features["class"])
    return image, mask, class_

autotune = tf.data.AUTOTUNE
create_tensorflow_record = True
DIR_TFRECORD = "/kaggle/working/blood_cell_cancer_with_mask"

if create_tensorflow_record == True:
    createTFRecord(pathfile=DIR_TFRECORD+"_train.tfrecord", dataset=df_train,
                   createTFRecord(pathfile=DIR_TFRECORD+"_val.tfrecord", dataset=df_val, size=1000),
                   createTFRecord(pathfile=DIR_TFRECORD+"_test.tfrecord", dataset=df_test, size=1000))

train_data = tf.data.TFRecordDataset(DIR_TFRECORD+"_train.tfrecord")
train_data = train_data.map(parse, num_parallel_calls=autotune)

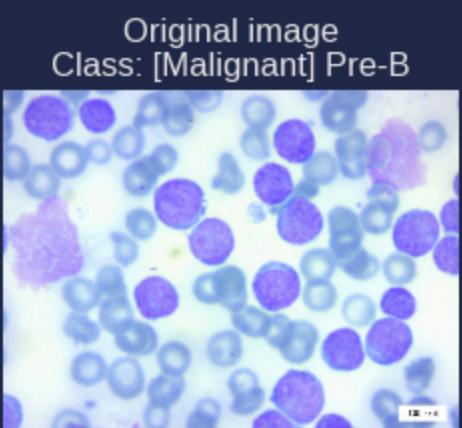
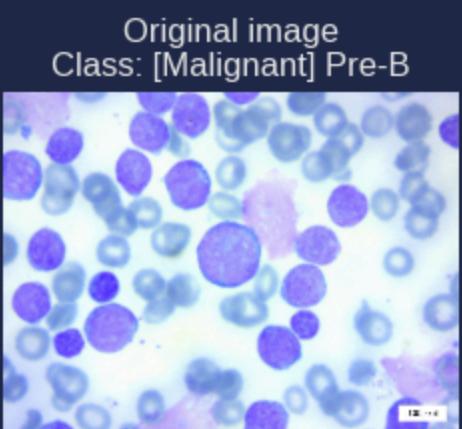
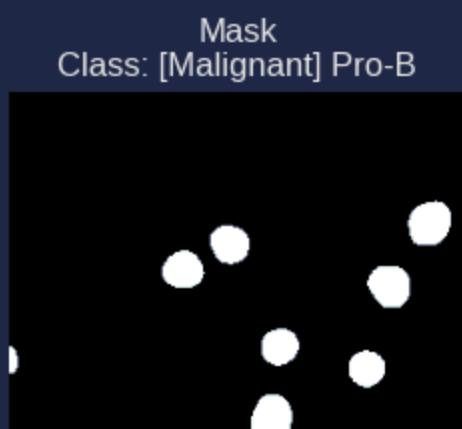
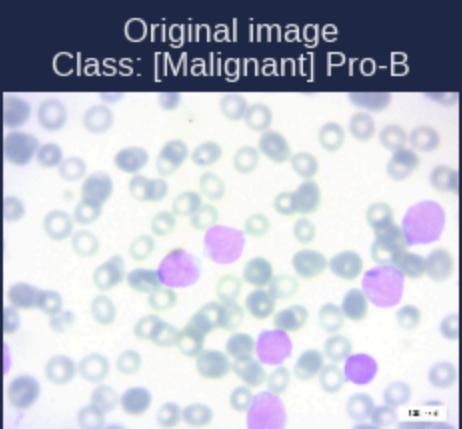
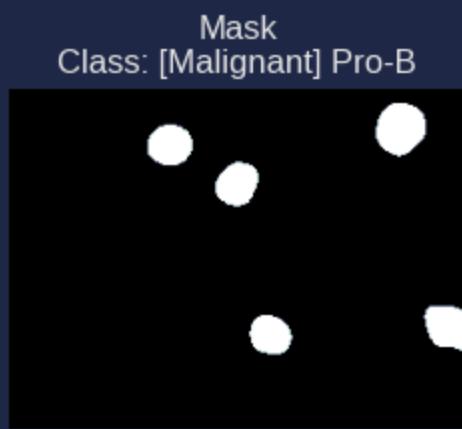
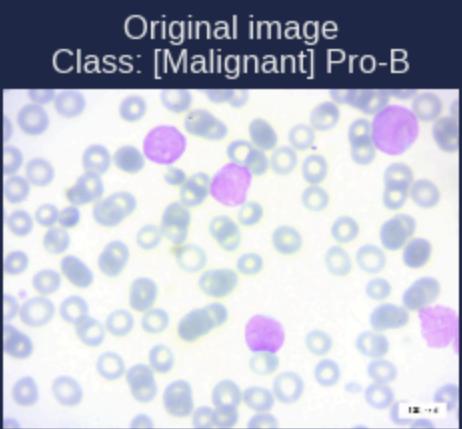
val_data = tf.data.TFRecordDataset(DIR_TFRECORD+"_val.tfrecord")
val_data = val_data.map(parse, num_parallel_calls=autotune)

test_data = tf.data.TFRecordDataset(DIR_TFRECORD+"_test.tfrecord")
test_data = test_data.map(parse, num_parallel_calls=autotune)

```

Let's take a preview

```
In [ ]: plt.subplots(4, 2, figsize=(5, 10))
for i, data in enumerate(train_data.take(4)):
    plt.subplot(4, 2, (i*2)+1)
    plt.imshow(data[0])
    plt.axis("off")
    plt.title(f"Original image\nClass: {classes[data[2][0]]}")
    plt.subplot(4, 2, (i*2)+2)
    plt.imshow(data[1], cmap="bone")
    plt.axis("off")
    plt.title(f"Mask\nClass: {classes[data[2][0]]}")
plt.tight_layout()
```



Loading the images for image classification

```
In [ ]: batch_size = 32

def classification_task(image, mask, class_):
    class_ohe = tf.one_hot(indices=tf.squeeze(class_, axis=0), depth=4)
    return tf.cast(image, dtype=tf.float32)/255, tf.cast(class_ohe, dtype=tf.

train_data_clf = train_data.map(classification_task, num_parallel_calls=autotune)
train_data_clf = train_data_clf.batch(batch_size)
train_data_clf = train_data_clf.prefetch(autotune)

val_data_clf = val_data.map(classification_task, num_parallel_calls=autotune)
val_data_clf = val_data_clf.batch(batch_size)
val_data_clf = val_data_clf.prefetch(autotune)

test_data_clf = test_data.map(classification_task, num_parallel_calls=autotune)
test_data_clf = test_data_clf.batch(batch_size)
test_data_clf = test_data_clf.prefetch(autotune)
```

Loading the images for image segmentation

```
In [ ]: def segmentation_task(image, mask, class_):
    return tf.cast(image, dtype=tf.float32)/255, tf.cast(mask, dtype=tf.float32)

train_data_sg = train_data.map(segmentation_task, num_parallel_calls=autotune)
train_data_sg = train_data_sg.batch(batch_size)
train_data_sg = train_data_sg.prefetch(autotune)

val_data_sg = val_data.map(segmentation_task, num_parallel_calls=autotune)
val_data_sg = val_data_sg.batch(batch_size)
val_data_sg = val_data_sg.prefetch(autotune)

test_data_sg = test_data.map(segmentation_task, num_parallel_calls=autotune)
test_data_sg = test_data_sg.batch(batch_size)
test_data_sg = test_data_sg.prefetch(autotune)
```

Loading the images for both tasks

```
In [ ]: def both_task(image, mask, class_):
    class_ohe = tf.one_hot(indices=tf.squeeze(class_, axis=0), depth=4)
    return tf.cast(image, dtype=tf.float32)/255, {"segmentation": tf.cast(mask, d

train_data_both = train_data.map(both_task, num_parallel_calls=autotune)
train_data_both = train_data_both.batch(batch_size)
train_data_both = train_data_both.prefetch(autotune)

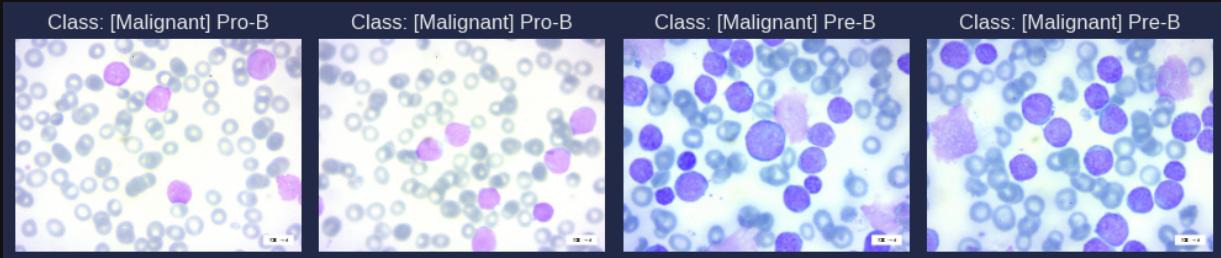
val_data_both = val_data.map(both_task, num_parallel_calls=autotune)
val_data_both = val_data_both.batch(batch_size)
val_data_both = val_data_both.prefetch(autotune)
```

```
test_data_both = test_data.map(both_task, num_parallel_calls=autotune)
test_data_both = test_data_both.batch(batch_size)
test_data_both = test_data_both.prefetch(autotune)
```

Training image visualization

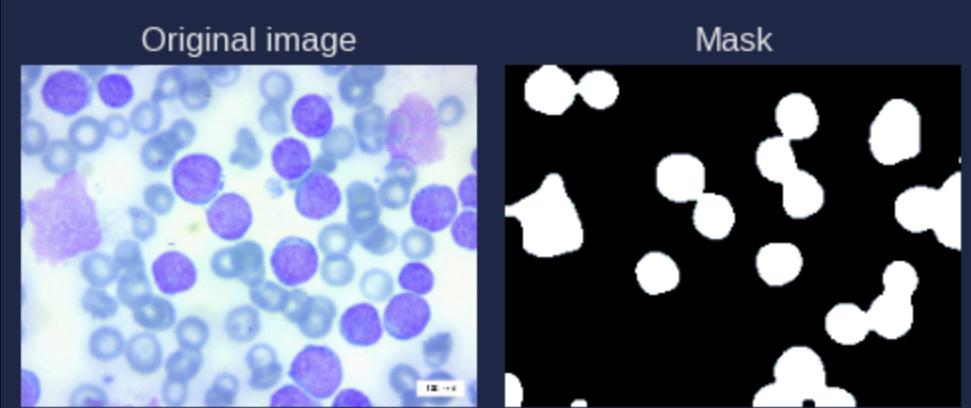
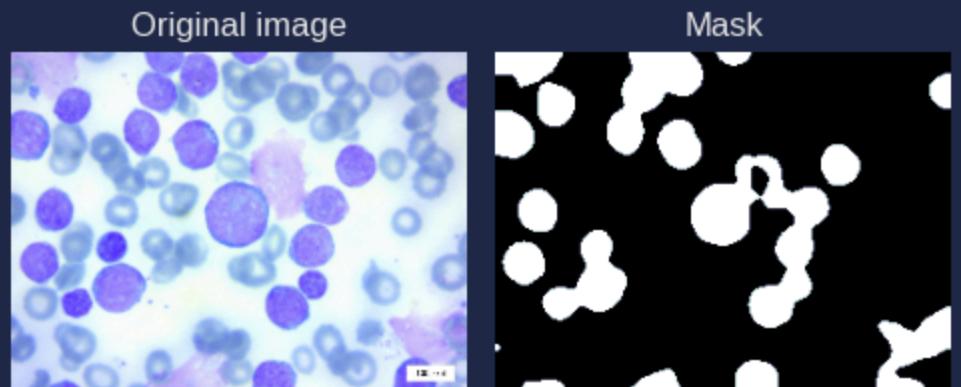
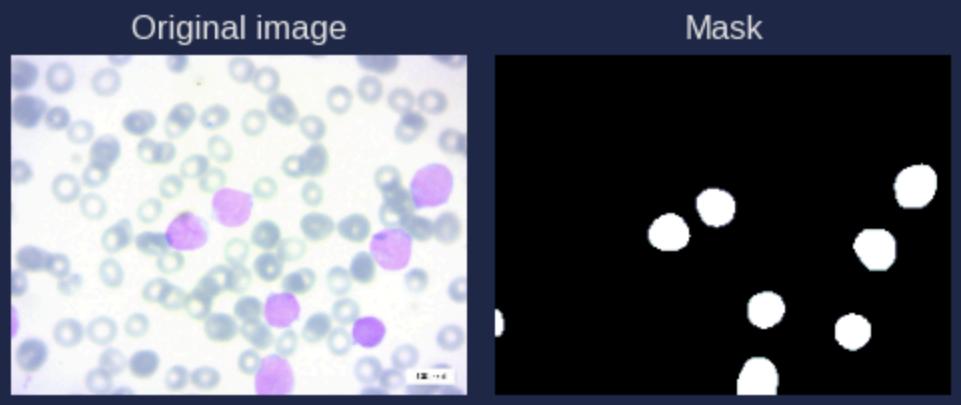
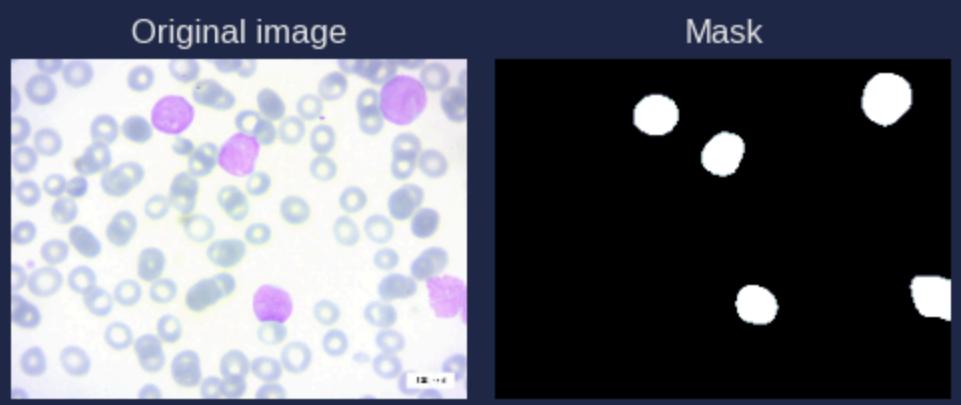
For classification

```
In [ ]: plt.subplots(1, 4, figsize=(10, 4))
for i, (image, class_ohe) in enumerate(train_data_clf.take(1).unbatch()):
    if i == 4:
        break
    plt.subplot(1, 4, i+1)
    plt.imshow(tf.cast(image*255, tf.uint8))
    plt.title(f"Class: {classes[np.argmax(class_ohe)]}")
    plt.axis("off")
plt.tight_layout()
```



For segmentation

```
In [ ]: plt.subplots(4, 2, figsize=(5, 10))
for i, (image, mask) in enumerate(train_data_sg.take(1).unbatch()):
    if i == 4:
        break
    plt.subplot(4, 2, (i*2)+1)
    plt.imshow(image)
    plt.axis("off")
    plt.title("Original image")
    plt.subplot(4, 2, (i*2)+2)
    plt.imshow(mask, cmap="bone")
    plt.axis("off")
    plt.title("Mask")
plt.tight_layout()
```



Handling imbalance image dataset

Calculating the class weights for the classes

```
In [ ]: df_train.type_cell.value_counts()
```

```
Out[ ]: type_cell
        [Malignant] early Pre-B    685
        [Malignant] Pre-B       668
        [Malignant] Pro-B      557
        Benign                  358
        Name: count, dtype: int64
```

```
In [ ]: import seaborn as sns
plt.figure(figsize=(15, 2))
sns.countplot(df_train, y="type_cell", alpha=1, order=df_train.type_cell.value_counts())
plt.ylabel("Type cell");
```



```
In [ ]: class_weight = dict(zip([i for i, _ in enumerate(classes)], compute_class_weight(class_weight)))
```

```
Out[ ]: {0: 0.8277372262773722,
         1: 1.0179533213644525,
         2: 0.8488023952095808,
         3: 1.5837988826815643}
```

Calculating the class weights for the masks

```
In [ ]: pixels_one = 0
pixels_zero = 0

for image, mask, class_ in train_data:
    aux = tf.reduce_sum(mask/255)
    pixels_one += aux
    pixels_zero += (mask.shape[0] * mask.shape[1]) - aux

total_pixels = pixels_one + pixels_zero
print(pixels_one, pixels_zero, total_pixels)
```

```
tf.Tensor(12867587.0, shape=(), dtype=float32) tf.Tensor(98609190.0, shape=(),
dtype=float32) tf.Tensor(111476776.0, shape=(), dtype=float32)
```

```
In [ ]: class_weight_pixel = {0:(total_pixels/(2*pixels_zero)).numpy(), 1:(total_pix
```

```
Out[ ]: {0: 0.56524533, 1: 4.3316894}
```

Diagnosing Acute Lymphoblastic Leukemia (ALL) by classifying and segmenting images of blood cells affected by cancer

Image classification

Creating a baseline model for image classification

```
In [ ]: def conv2d(filters):
    return Conv2D(filters, kernel_size=(2, 2), padding='same', strides=(1, 1))

def block_conv(inp, filters=128, dropout=0.4):
    x = conv2d(filters)(inp)
    x = ReLU()(x)
    x = BatchNormalization()(x)

    x = conv2d(filters//2)(x)
    x = ReLU()(x)
    x = BatchNormalization()(x)

    x = MaxPooling2D(pool_size=(2, 2))(x)
    return Dropout(dropout)(x)

def classification_base_architecture(inp, name_output=None):
    x = block_conv(inp, filters=128, dropout=0.4)
    x = block_conv(x, filters=128, dropout=0.4)

    x = GlobalAveragePooling2D()(x)
    x = Flatten()(x)
    x = Dense(32, activation="relu")(x)
    x = Dense(16, activation="relu")(x)
    x = Dense(len(classes), activation="softmax", name=name_output)(x)
    return x

def get_model():
    tf.keras.backend.clear_session()

    inp = Input(shape=img_shape+(3,))
    x = classification_base_architecture(inp)
    model = Model(inputs=inp, outputs=x)

    model.compile(loss=CategoricalCrossentropy(), optimizer=Adam(learning_ra
```

```
    return model  
  
model_base = get_model()
```

```
In [ ]: model_base.summary()
```

Model: "functional"

Layer (type)	Output Shape
input_layer (InputLayer)	(None, 192, 256, 3)
conv2d (Conv2D)	(None, 192, 256, 128)
re_lu (ReLU)	(None, 192, 256, 128)
batch_normalization (BatchNormalization)	(None, 192, 256, 128)
conv2d_1 (Conv2D)	(None, 192, 256, 64)
re_lu_1 (ReLU)	(None, 192, 256, 64)
batch_normalization_1 (BatchNormalization)	(None, 192, 256, 64)
max_pooling2d (MaxPooling2D)	(None, 96, 128, 64)
dropout (Dropout)	(None, 96, 128, 64)
conv2d_2 (Conv2D)	(None, 96, 128, 128)
re_lu_2 (ReLU)	(None, 96, 128, 128)
batch_normalization_2 (BatchNormalization)	(None, 96, 128, 128)
conv2d_3 (Conv2D)	(None, 96, 128, 64)
re_lu_3 (ReLU)	(None, 96, 128, 64)
batch_normalization_3 (BatchNormalization)	(None, 96, 128, 64)
max_pooling2d_1 (MaxPooling2D)	(None, 48, 64, 64)
dropout_1 (Dropout)	(None, 48, 64, 64)
global_average_pooling2d (GlobalAveragePooling2D)	(None, 64)
flatten (Flatten)	(None, 64)
dense (Dense)	(None, 32)
dense_1 (Dense)	(None, 16)
dense_2 (Dense)	(None, 4)

Total params: 104,436 (407.95 KB)

Trainable params: 103,668 (404.95 KB)

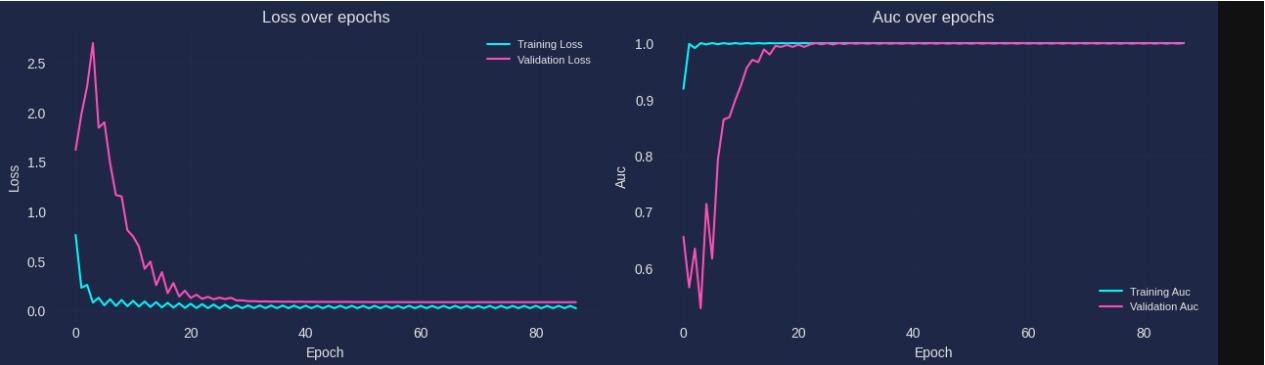
Non-trainable params: 768 (3.00 KB)

Let's train

```
In [ ]: def train_model(model, train_data, val_data, epochs=100, version="base"):  
    checkpoint_filepath = f'/kaggle/working/models/model_{version}.keras'  
    checkpoint = ModelCheckpoint(checkpoint_filepath, monitor='val_loss', save  
    early = EarlyStopping(monitor='val_loss', patience=8, verbose=1, restore  
    reduce = ReduceLROnPlateau(monitor='val_loss', patience=3)  
    csvlogger = CSVLogger(f"/kaggle/working/histories/history_model_{version}  
  
        history = model.fit(train_data,  
                            batch_size=batch_size,  
                            epochs=epochs,  
                            callbacks=[early,  
                                      reduce,  
                                      checkpoint,  
                                      csvlogger],  
                            validation_data=val_data,  
                            class_weight=class_weight,  
                            steps_per_epoch=len(df_train)//batch_size,  
                            validation_steps=len(df_val)//batch_size,  
                            validation_batch_size=batch_size  
        )  
    return history
```

```
In [ ]: history_base = train_model(model_base, train_data_clf, val_data_clf, epochs=
```

```
In [ ]: def delete_empty_subplots(fig, axs):  
    for ax_row in axs:  
        for ax in ax_row:  
            if ax.title.get_text() == "":  
                fig.delaxes(ax)  
    return fig, axs  
  
def show_history(version, parameters=None):  
    df_hist = pd.read_csv(f"/kaggle/working/histories/history_model_{version}")  
    fig, axs = plt.subplots(2, 4, figsize=(25, 7))  
  
    for i, col in enumerate(parameters):  
        plt.subplot(2, 4, i+1)  
        param = col[0].replace('_', ' ').title()  
        plt.plot(df_hist.index, df_hist[col[0]], label=f"Training {param}")  
        plt.plot(df_hist.index, df_hist[col[1]], label=f"Validation {param}")  
        plt.ylabel(col[0].capitalize().replace("_", " "))  
        plt.xlabel("Epoch")  
        plt.legend(fontsize=8)  
        plt.grid(True, alpha=0.3, which="both")  
        plt.title(f"{param} over epochs")  
  
    fig, axs = delete_empty_subplots(fig, axs)  
    plt.tight_layout()  
    plt.show()  
    return  
  
show_history(version="base", parameters=[["loss", "val_loss"], ["auc", "val_a
```



Evaluating the baseline model for image classification

```
In [ ]: model_base.evaluate(test_data_clf)
16/16 ━━━━━━━━━━━━━━━━ 1s 35ms/step - auc: 0.9999 - loss: 0.0631
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()` function when building your dataset.
    self.gen.throw(typ, value, traceback)
```

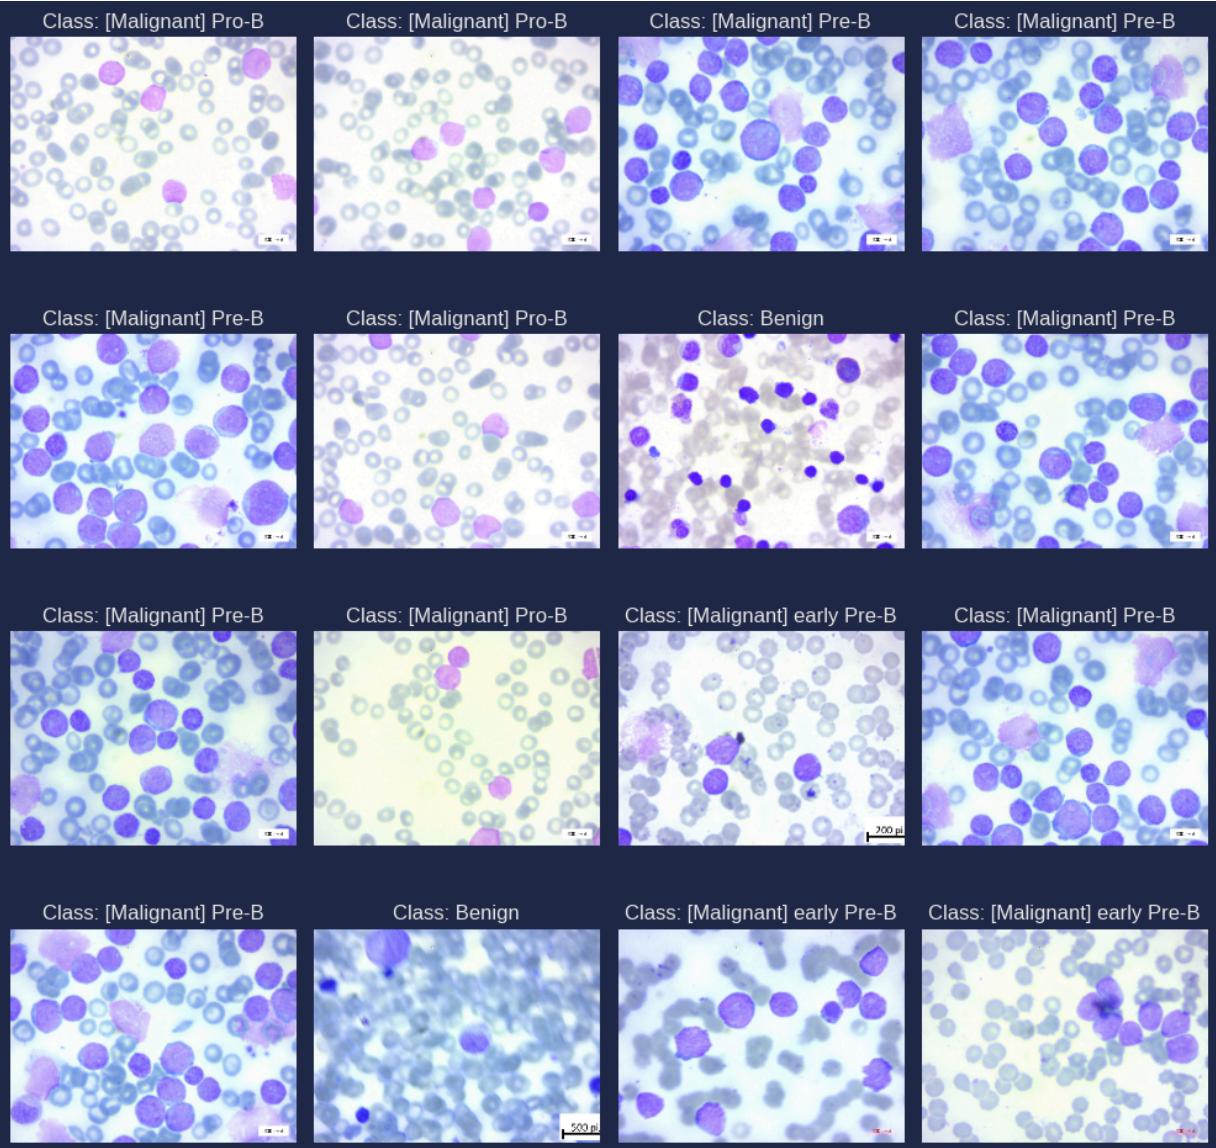
Out []: [0.06927863508462906, 0.9997674822807312]

Applying data augmentation technique on images

```
In [ ]: def data_augmentation(image, mask, class_):
    img = random_flip_left_right(image, seed=seed_value)
    img = random_flip_up_down(img, seed=seed_value)
    img = random_brightness(img, max_delta=0.01, seed=seed_value)
    img = random_contrast(img, lower=0.5, upper=1, seed=seed_value)
    img = random_saturation(img, lower=0.5, upper=1, seed=seed_value)
    class_ohe = tf.one_hot(indices=tf.squeeze(class_, axis=0), depth=4)
    return tf.cast(image, dtype=tf.float32)/255, tf.cast(class_ohe, dtype=tf.int32)

train_data_clf_aug = train_data.map(data_augmentation, num_parallel_calls=auto)
train_data_clf_aug = train_data_clf_aug.batch(batch_size)
train_data_clf_aug = train_data_clf_aug.prefetch(autotune)

plt.subplots(4, 4, figsize=(10, 10))
for i, (image, class_ohe) in enumerate(train_data_clf_aug.take(1).unbatch()):
    if i == 16:
        break
    plt.subplot(4, 4, i+1)
    plt.imshow(tf.cast(image*255, tf.uint8))
    plt.title(f"Class: {classes[np.argmax(class_ohe)]}")
    plt.axis("off")
plt.tight_layout()
```

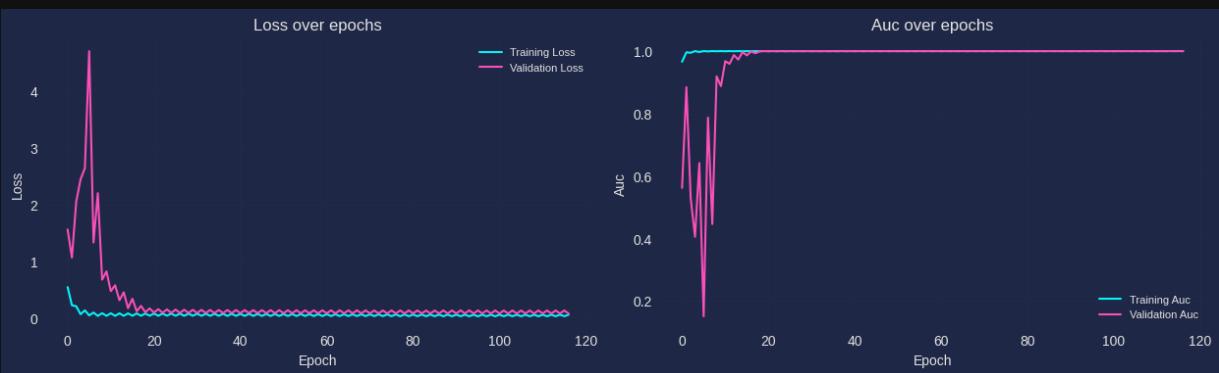


Training a baseline model with data augmentation

```
In [ ]: model_base_aug = get_model()
history_base_aug = train_model(model_base_aug, train_data_clf_aug, val_data)
```



```
In [ ]: show_history(version="base_aug", parameters=[["loss", "val_loss"], ["auc", "v
```



Evaluating the baseline model with data augmentation

```
In [ ]: model_base_aug.evaluate(test_data_clf)
```

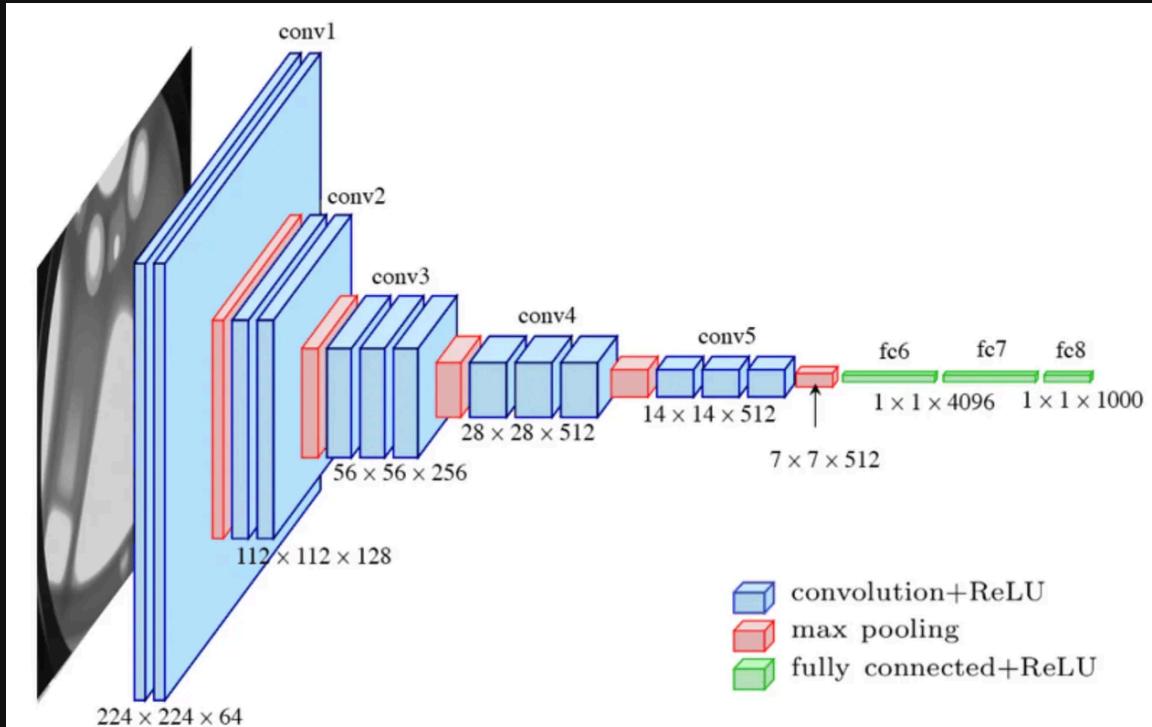
```
16/16 ━━━━━━━━━━━━━━━━ 1s 34ms/step - auc: 0.9995 - loss: 0.0655
```

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()` function when building your dataset.
```

```
    self.gen.throw(typ, value, traceback)
```

```
Out[ ]: [0.07973799854516983, 0.9987912774085999]
```

Transfer learning using the architecture proposed by the Visual Geometry Group (VGG)



```
In [ ]: def get_backbone():
    backbone = VGG16(include_top=False, weights='imagenet', input_shape=img)
    backbone.trainable = False
    return backbone

def classification_w_backbone_architecture(inp, name_output=None):
    backbone = get_backbone()
    x = backbone(inp, training=False)
    x = Flatten()(x)
    x = Dense(32, activation="relu")(x)
    x = Dense(16, activation="relu")(x)
    x = Dense(len(classes), activation="softmax", name=name_output)(x)
```

```

    return x

def get_model_w_backbone():
    tf.keras.backend.clear_session()

    inp = Input(shape=img_shape+(3,))
    x = classification_w_backbone_architecture(inp)
    model = Model(inputs=inp, outputs=x)

    model.compile(loss=CategoricalCrossentropy(), optimizer=Adam(learning_rate))
    return model

model_backbone = get_model_w_backbone()

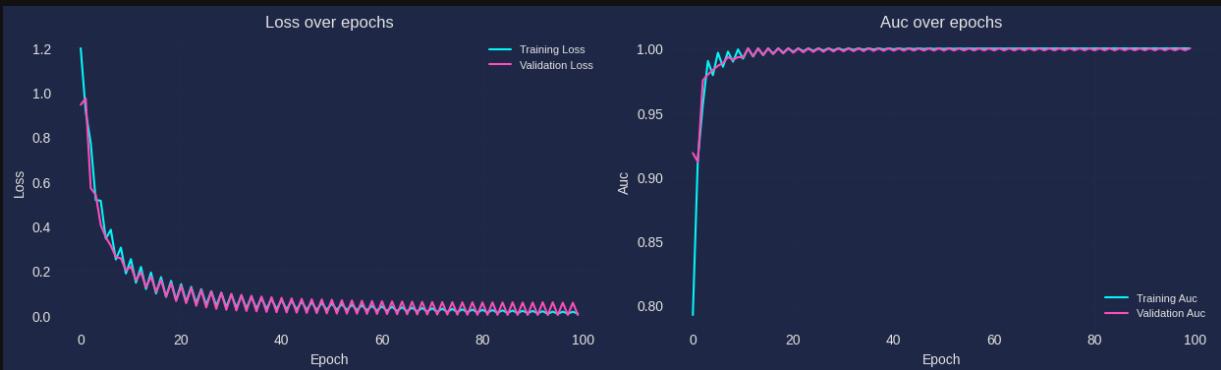
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 ━━━━━━━━━━━━━━━━ 2s 0us/step

```

Let's train

```
In [ ]: history_backbone = train_model(model_backbone, train_data_clf_aug, val_data_
```

```
In [ ]: show_history(version="backbone", parameters=[["loss", "val_loss"], ["auc", "v
```



Evaluating the model with a VGG backbone

```
In [ ]: model_backbone.evaluate(test_data_clf)
```

```
16/16 ━━━━━━━━━━━━━━━━ 1s 58ms/step - auc: 0.9996 - loss: 0.0469
```

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()` function when building your dataset.
  self.gen.throw(typ, value, traceback)
```

```
Out[ ]: [0.051316503435373306, 0.9995214939117432]
```

Evaluation metrics for classification task

```
In [ ]: def get_scores(y_true, y_pred, y_pred_proba):
    return {'AUC_ROC':roc_auc_score(y_true, y_pred_proba, multi_class='ovr'),
            'F1_Score':f1_score(y_true, y_pred, average='weighted'),
            'Accuracy':accuracy_score(y_true, y_pred),
            'Precision':precision_score(y_true, y_pred, average='weighted'),
            'Recall':recall_score(y_true, y_pred, average='weighted')}

for i, (img, class_) in enumerate(test_data_clf):
    if i == 0:
        y_true = class_.numpy()
    else:
        y_true = np.concatenate((y_true, class_.numpy()), axis=0)

scores = []
for model in [model_base, model_base_aug, model_backbone]:
    y_test_pred_proba = model.predict(test_data_clf, verbose=0)
    y_test_pred = tf.math.round(y_test_pred_proba)
    scores.append(get_scores(y_true, y_test_pred, y_test_pred_proba))

comparison = pd.DataFrame(data=scores, index=["Baseline Model", "Baseline Mo
comparison.style.highlight_max(color = 'green', axis = 0).highlight_min(colo
```

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of dat
a; interrupting training. Make sure that your dataset or generator can genera
te at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()`` function when building your dataset.
    self.gen.throw(typ, value, traceback)
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of dat
a; interrupting training. Make sure that your dataset or generator can genera
te at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()`` function when building your dataset.
    self.gen.throw(typ, value, traceback)
```

Out []:

	AUC_ROC	F1_Score	Accuracy	Precision	Recall
Baseline Model	0.999683	0.985588	0.983573	0.987708	0.983573
Baseline Model with DA	0.999239	0.983538	0.981520	0.985767	0.981520
Model with VGG16	0.999667	0.979537	0.979466	0.980103	0.979466

Image segmentation

Creating a baseline model for image segmentation

In []:

```
def encoder(inp, filters):
    x1 = block_conv(inp, filters=filters, dropout=0.4)
    x2 = block_conv(x1, filters=filters, dropout=0.4)
    return x1, x2

def bottleneck(inp, filters):
```

```

x = conv2d(filters)(inp)
x = ReLU()(x)
x = BatchNormalization()(x)

x = conv2d(filters)(inp)
x = ReLU()(x)
return BatchNormalization()(x)

def conv2dtranspose(filters):
    return Conv2DTranspose(filters=filters, activation="relu", kernel_size=2, strides=2)

def up_block(inp, output_enc, filters, dropout=0.4):
    x = conv2d(filters)(inp)
    x = Concatenate()([inp, output_enc])
    x = ReLU()(x)
    x = BatchNormalization()(x)

    x = conv2d(filters)(inp)
    x = ReLU()(x)
    x = BatchNormalization()(x)

    x = conv2dtranspose(filters)(x)
    return Dropout(dropout)(x)

def decoder(inp, ouput_block1_enc, ouput_block2_enc, filters):
    x1 = up_block(inp, ouput_block2_enc, filters, dropout=0.4)
    x2 = up_block(x1, ouput_block1_enc, filters, dropout=0.4)
    return x2

def segmentation_architecture(inp, name_output=None):
    ouput_block1_enc, ouput_block2_enc = encoder(inp, filters=128)
    x = bottleneck(ouput_block2_enc, filters=128)
    x = decoder(x, ouput_block1_enc, ouput_block2_enc, filters=128)
    x = Conv2D(filters=1, activation="sigmoid", kernel_size=(2, 2), padding='same')
    return x

def get_model_segmentation():
    tf.keras.backend.clear_session()

    inp = Input(shape=img_shape+(3,))
    x = segmentation_architecture(inp)
    model = Model(inputs=inp, outputs=x)

    model.compile(loss=Dice(),
                  optimizer=Adam(learning_rate=1e-3),
                  metrics=[BinaryIoU(), BinaryAccuracy()])
    return model

model_base_sg = get_model_segmentation()

```

In []: model_base_sg.summary()

Model: "functional"

Layer (type)	Output Shape
input_layer (InputLayer)	(None, 192, 256, 3)
conv2d (Conv2D)	(None, 192, 256, 128)
re_lu (ReLU)	(None, 192, 256, 128)
batch_normalization (BatchNormalization)	(None, 192, 256, 128)
conv2d_1 (Conv2D)	(None, 192, 256, 64)
re_lu_1 (ReLU)	(None, 192, 256, 64)
batch_normalization_1 (BatchNormalization)	(None, 192, 256, 64)
max_pooling2d (MaxPooling2D)	(None, 96, 128, 64)
dropout (Dropout)	(None, 96, 128, 64)
conv2d_2 (Conv2D)	(None, 96, 128, 128)
re_lu_2 (ReLU)	(None, 96, 128, 128)
batch_normalization_2 (BatchNormalization)	(None, 96, 128, 128)
conv2d_3 (Conv2D)	(None, 96, 128, 64)
re_lu_3 (ReLU)	(None, 96, 128, 64)
batch_normalization_3 (BatchNormalization)	(None, 96, 128, 64)
max_pooling2d_1 (MaxPooling2D)	(None, 48, 64, 64)
dropout_1 (Dropout)	(None, 48, 64, 64)
conv2d_5 (Conv2D)	(None, 48, 64, 128)
re_lu_5 (ReLU)	(None, 48, 64, 128)
batch_normalization_5 (BatchNormalization)	(None, 48, 64, 128)
conv2d_7 (Conv2D)	(None, 48, 64, 128)
re_lu_7 (ReLU)	(None, 48, 64, 128)
batch_normalization_7 (BatchNormalization)	(None, 48, 64, 128)
conv2d_transpose (Conv2DTranspose)	(None, 96, 128, 128)
dropout_2 (Dropout)	(None, 96, 128, 128)

conv2d_9 (Conv2D)	(None, 96, 128, 128)
re_lu_9 (ReLU)	(None, 96, 128, 128)
batch_normalization_9 (BatchNormalization)	(None, 96, 128, 128)
conv2d_transpose_1 (Conv2DTranspose)	(None, 192, 256, 128)
dropout_3 (Dropout)	(None, 192, 256, 128)
conv2d_10 (Conv2D)	(None, 192, 256, 1)

Total params: 399,361 (1.52 MB)

Trainable params: 397,825 (1.52 MB)

Non-trainable params: 1,536 (6.00 KB)

```
In [ ]: def show_predictions(dataset=None):
    for image, _ in dataset.shuffle(100).take(1):
        pred = model_base_sg.predict(image)
        pred_mask = pred[0] >= 0.5
        plt.subplots(1, 2, figsize=(6, 3))

        plt.subplot(1, 2, 1)
        plt.imshow(image[0])
        plt.axis("off")
        plt.title(f"Original image")

        plt.subplot(1, 2, 2)
        plt.imshow(pred_mask, cmap="bone")
        plt.axis("off")
        plt.title(f"Mask")

    plt.tight_layout()
    plt.show()
    break
return

class DisplayCallback(Callback):
    def on_epoch_end(self, epoch, logs=None):
        clear_output(wait=True)
        show_predictions(dataset=val_data_sg)
        print ('\nSample Prediction after epoch {} \n'.format(epoch+1))
    return
```

Let's train

```
In [ ]: def train_model_segmentation(model, train_data, val_data, epochs=100, version=1):
    checkpoint_filepath = f'/kaggle/working/models/model_{version}.keras'
    checkpoint = ModelCheckpoint(checkpoint_filepath, monitor='val_loss', save_best_only=True)
    early = EarlyStopping(monitor='val_loss', patience=8, verbose=1, restore_best_weights=True)
    reduce = ReduceLROnPlateau(monitor='val_loss', patience=3)
    csvlogger = CSVLogger(f'/kaggle/working/histories/history_model_{version}.csv')
```

```

        history = model.fit(train_data,
                             batch_size=batch_size,
                             epochs=epochs,
                             callbacks=[early,
                                        reduce,
                                        checkpoint,
                                        csvlogger,
                                        DisplayCallback()],
                             validation_data=val_data,
                             class_weight=class_weight_pixel,
                             steps_per_epoch=len(df_train)//batch_size,
                             validation_steps=len(df_val)//batch_size,
                             validation_batch_size=batch_size
                            )
    return history

```

In []: history_base_sg = train_model_segmentation(model_base_sg, train_data_sg, val

1/1 ————— 0s 27ms/step

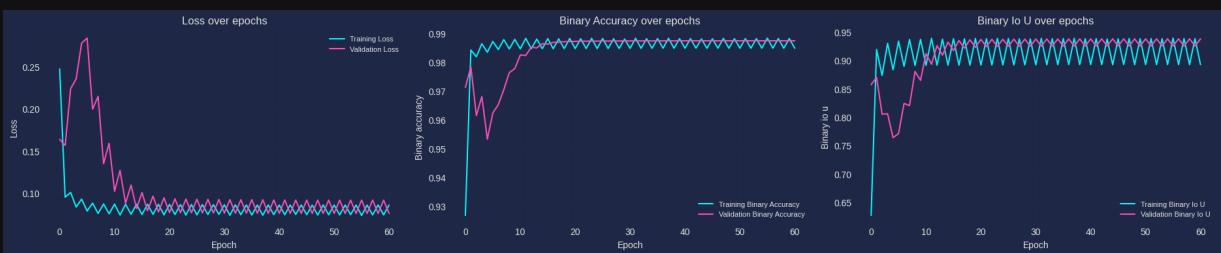


Sample Prediction after epoch 61

70/70 ————— 19s 269ms/step - binary_accuracy: 0.9849 - binary_io_u: 0.9139 - loss: 0.0848 - val_binary_accuracy: 0.9873 - val_binary_io_u: 0.9375 - val_loss: 0.0755 - learning_rate: 1.0000e-13
Epoch 61: early stopping

Restoring model weights from the end of the best epoch: 53.

In []: show_history(version="base_segmentation", parameters=[["loss", "val_loss"], ["binary_accuracy", "val_binary_accuracy"], ["binary_io_u", "val_binary_io_u"]])



Evaluating the baseline model for segmentation

```
In [ ]: result_test_sg = model_base_sg.evaluate(test_data_sg)
result_test_sg
16/16 ━━━━━━━━━━━━ 1s 83ms/step - binary_accuracy: 0.98
51 - binary_io_u: 0.9295 - loss: 0.0845
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()` function when building your dataset.
    self.gen.throw(typ, value, traceback)

Out[ ]: [0.08604717999696732, 0.9265201091766357, 0.9849867820739746]
```

Applying data augmentation technique on images and masks

```
In [ ]: def data_transformation(image, mask, class_):
    img = random_brightness(image, max_delta=0.01, seed=seed_value)
    img = random_contrast(img, lower=0.5, upper=1, seed=seed_value)
    img = random_saturation(img, lower=0.5, upper=1, seed=seed_value)
    aux = tf.concat([img, mask], -1)

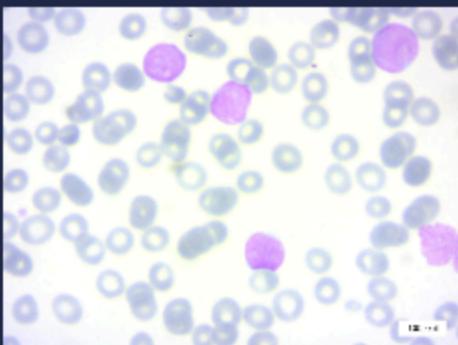
    aux_ = random_flip_left_right(aux, seed=seed_value)
    aux_ = random_flip_up_down(aux_, seed=seed_value)

    return tf.cast(aux_[..., :3], tf.float32)/255, tf.cast(aux_[..., 3:], tf.

train_data_sg_aug = train_data.map(data_transformation, num_parallel_calls=a
train_data_sg_aug = train_data_sg_aug.batch(batch_size)
train_data_sg_aug = train_data_sg_aug.prefetch(autotune)
```

```
In [ ]: plt.subplots(4, 2, figsize=(5, 10))
for i, (image, mask) in enumerate(train_data_sg_aug.take(1).unbatch()):
    if i == 4:
        break
    plt.subplot(4, 2, (i*2)+1)
    plt.imshow(image)
    plt.axis("off")
    plt.title(f"Original image")
    plt.subplot(4, 2, (i*2)+2)
    plt.imshow(mask, cmap="bone")
    plt.axis("off")
    plt.title(f"Mask")
plt.tight_layout()
```

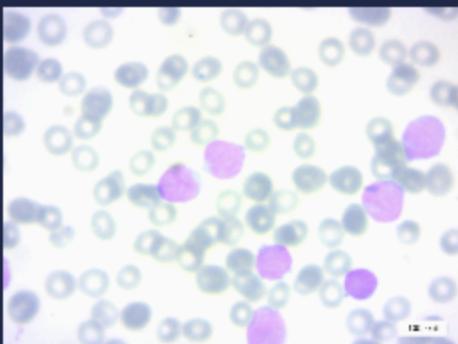
Original image



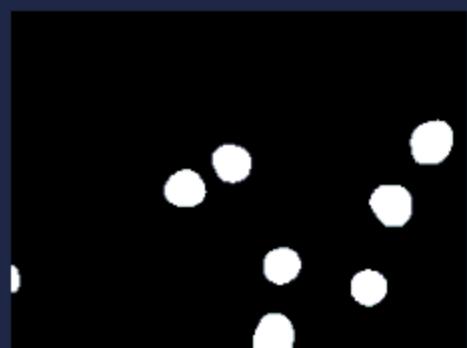
Mask



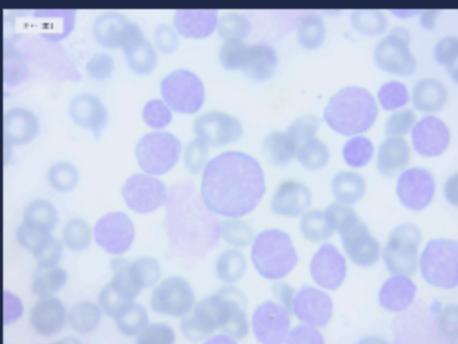
Original image



Mask



Original image



Mask



Original image



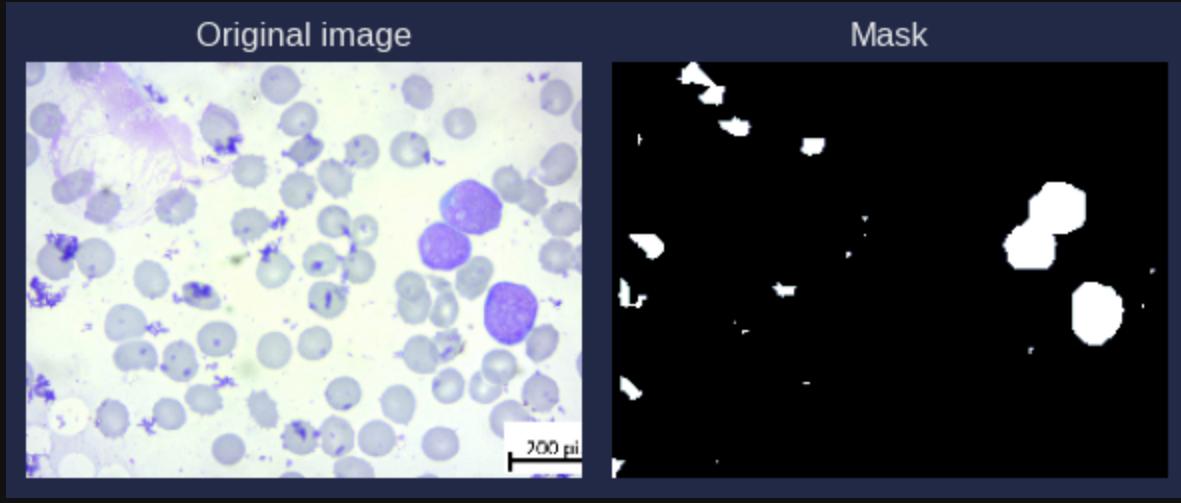
Mask



Training a baseline model with data augmentation

```
In [ ]: model_base_sg_aug = get_model_segmentation()  
history_base_sg_aug = train_model_segmentation(model_base_sg_aug, train_data)
```

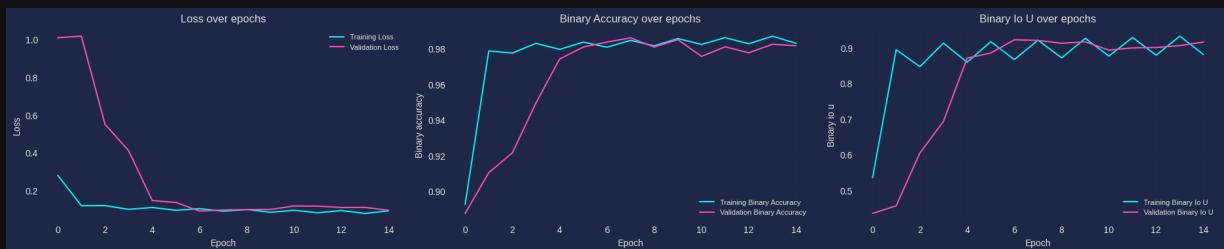
1/1 ————— 0s 70ms/step



Sample Prediction after epoch 15

```
70/70 ————— 19s 270ms/step - binary_accuracy: 0.  
9830 - binary_io_u: 0.9041 - loss: 0.0910 - val_binary_accuracy: 0.9816 - val  
_binary_io_u: 0.9156 - val_loss: 0.0941 - learning_rate: 1.0000e-05  
Epoch 15: early stopping  
Restoring model weights from the end of the best epoch: 7.
```

```
In [ ]: show_history(version="base_segmentation_aug", parameters=[["loss", "val_loss",  
["binary_accuracy", "v  
["binary_io_u", "val_b
```



Evaluating the baseline model for image segmentation with data augmentation

```
In [ ]: result_test_sg_aug = model_base_sg_aug.evaluate(test_data_sg)  
result_test_sg_aug
```

```
16/16 ————— 1s 82ms/step - binary_accuracy: 0.98  
14 - binary_io_u: 0.9150 - loss: 0.0982
```

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()` function when building your dataset.
    self.gen.throw(typ, value, traceback)
```

```
Out[ ]: [0.09920085221529007, 0.9123088121414185, 0.9814990758895874]
```

Evaluation metrics for segmentation task

```
In [ ]: for i, (img, mask) in enumerate(test_data_sg):
    aux = tf.reshape(mask, [-1])
    if i == 0:
        y_true = aux.numpy()
    else:
        y_true = np.concatenate((y_true, aux.numpy()), axis=0)

scores = []
for model in [model_base_sg, model_base_sg_aug]:
    y_test_pred_proba = model.predict(test_data_sg, verbose=0)
    y_test_pred_proba = tf.reshape(y_test_pred_proba, [-1])
    y_test_pred = tf.math.round(y_test_pred_proba)
    result_test = model.evaluate(test_data_sg)
    score = get_scores(y_true, y_test_pred, y_test_pred_proba)
    score["BinaryIOU"] = result_test[1]
    score["BinaryAccuracy"] = result_test[2]
    scores.append(score)

comparison = pd.DataFrame(data=scores, index=["Baseline Model", "Baseline Model with DA"])
comparison.style.highlight_max(color = 'green', axis = 0).highlight_min(color = 'red')
```

```
16/16 ━━━━━━━━━━━━━━━━ 1s 82ms/step - binary_accuracy: 0.98
51 - binary_io_u: 0.9295 - loss: 0.0845
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()` function when building your dataset.
    self.gen.throw(typ, value, traceback)
```

```
16/16 ━━━━━━━━━━━━━━━━ 1s 82ms/step - binary_accuracy: 0.98
14 - binary_io_u: 0.9150 - loss: 0.0982
```

	AUC_ROC	F1_Score	Accuracy	Precision	Recall	BinaryIOU	BinaryAccuracy
Baseline Model	0.987874	0.984820	0.984987	0.984809	0.984987	0.926520	0.984987
Baseline Model with DA	0.989966	0.981490	0.981499	0.981481	0.981499	0.912309	0.981499

Image segmentation and classification

Creating a baseline model for image segmentation and classification

```
In [ ]: def get_model_both():
    tf.keras.backend.clear_session()

    inp = Input(shape=img_shape+(3,))
    seg_output = segmentation_architecture(inp, name_output="segmentation")
    clf_output = classification_w_backbone_architecture(inp, name_output="cla

    model = Model(inputs=inp, outputs={"segmentation":seg_output,
                                       "classification":clf_output})

    model.compile(loss={"segmentation":Dice(),
                        "classification":CategoricalCrossentropy()},
                  optimizer=Adam(learning_rate=1e-3),
                  metrics={"segmentation": [BinaryIoU(), BinaryAccuracy()],
                           "classification": [AUC()]})
    )
    return model

model_base_both = get_model_both()
```

```
In [ ]: model_base_both.summary()
```

Model: "functional"

Layer (type)	Output Shape	Params
input_layer (InputLayer)	(None, 192, 256, 3)	0
conv2d (Conv2D)	(None, 192, 256, 128)	1,664
re_lu (ReLU)	(None, 192, 256, 128)	0
batch_normalization (BatchNormalization)	(None, 192, 256, 128)	512
conv2d_1 (Conv2D)	(None, 192, 256, 64)	32,832
re_lu_1 (ReLU)	(None, 192, 256, 64)	0
batch_normalization_1 (BatchNormalization)	(None, 192, 256, 64)	256
max_pooling2d (MaxPooling2D)	(None, 96, 128, 64)	0
dropout (Dropout)	(None, 96, 128, 64)	0
conv2d_2 (Conv2D)	(None, 96, 128, 128)	32,896
re_lu_2 (ReLU)	(None, 96, 128, 128)	0
batch_normalization_2 (BatchNormalization)	(None, 96, 128, 128)	512
conv2d_3 (Conv2D)	(None, 96, 128, 64)	32,832
re_lu_3 (ReLU)	(None, 96, 128, 64)	0
batch_normalization_3 (BatchNormalization)	(None, 96, 128, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 48, 64, 64)	0
dropout_1 (Dropout)	(None, 48, 64, 64)	0
conv2d_5 (Conv2D)	(None, 48, 64, 128)	32,896
re_lu_5 (ReLU)	(None, 48, 64, 128)	0
batch_normalization_5 (BatchNormalization)	(None, 48, 64, 128)	512
conv2d_7 (Conv2D)	(None, 48, 64, 128)	65,664
re_lu_7 (ReLU)	(None, 48, 64, 128)	0
batch_normalization_7 (BatchNormalization)	(None, 48, 64, 128)	512
conv2d_transpose (Conv2DTranspose)	(None, 96, 128, 128)	65,664

dropout_2 (Dropout)	(None, 96, 128, 128)	0
conv2d_9 (Conv2D)	(None, 96, 128, 128)	65,664
vgg16 (Functional)	(None, 512)	14,714,688
re_lu_9 (ReLU)	(None, 96, 128, 128)	0
flatten (Flatten)	(None, 512)	0
batch_normalization_9 (BatchNormalization)	(None, 96, 128, 128)	512
dense (Dense)	(None, 32)	16,416
conv2d_transpose_1 (Conv2DTranspose)	(None, 192, 256, 128)	65,664
dense_1 (Dense)	(None, 16)	528
dropout_3 (Dropout)	(None, 192, 256, 128)	0
classification (Dense)	(None, 4)	68
segmentation (Conv2D)	(None, 192, 256, 1)	513

Total params: 15,131,061 (57.72 MB)

Trainable params: 414,837 (1.58 MB)

Non-trainable params: 14,716,224 (56.14 MB)

Let's plot the model

```
In [ ]: plot_model(model_base_both, expand_nested=True, dpi=50, show_shapes=True)
```

Function for display some predictions during the training

```
In [ ]: def show_predictions(dataset=None):
    for image, _ in dataset.shuffle(100).take(1):
        pred = model_base_both.predict(image)
        pred_mask, pred_class = pred["segmentation"][0], pred["classification"]
        pred_mask = pred_mask >= 0.5

        plt.subplots(1, 2, figsize=(6, 3))

        plt.subplot(1, 2, 1)
        plt.imshow(image[0])
        plt.axis("off")
        plt.title(f"Original image\nClass: {classes[np.argmax(pred_class)]}")

        plt.subplot(1, 2, 2)
        plt.imshow(pred_mask, cmap="bone")
        plt.axis("off")
        plt.title(f"Mask\nClass: {classes[np.argmax(pred_class)]}")
```

```
    plt.tight_layout()
    plt.show()
    break
return

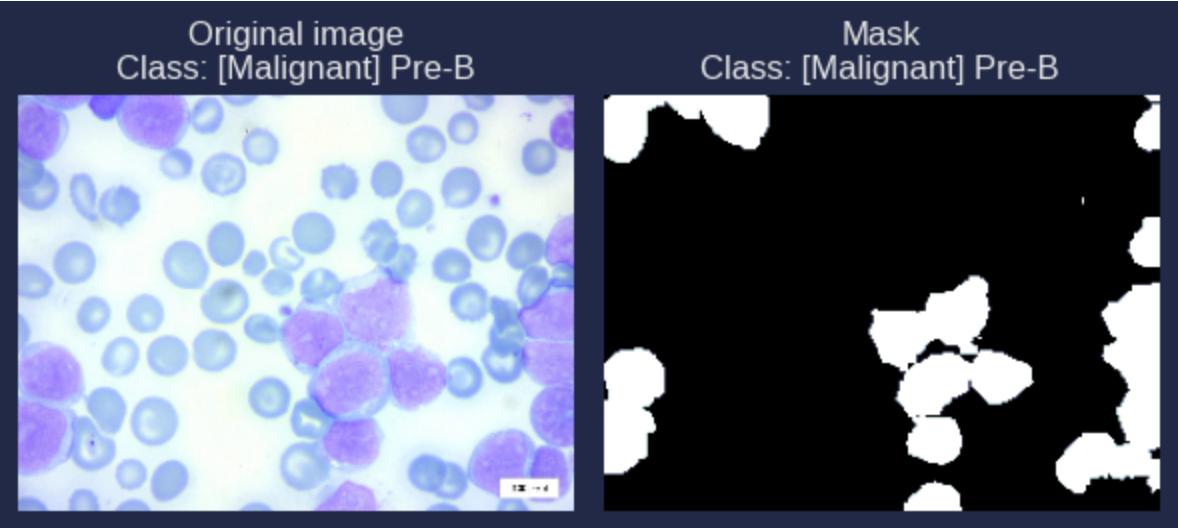
class DisplayCallback(Callback):
    def on_epoch_end(self, epoch, logs=None):
        clear_output(wait=True)
        show_predictions(dataset=val_data_sg)
        print ('\nSample Prediction after epoch {} \n'.format(epoch+1))
        return
```

Let's train

```
In [ ]: def train_model_both(model, train_data, val_data, epochs=100, version="base"):
    checkpoint_filepath = f'/kaggle/working/models/model_{version}.keras'
    checkpoint = ModelCheckpoint(checkpoint_filepath, monitor='val_loss', save_best_only=True)
    early = EarlyStopping(monitor='val_loss', patience=8, verbose=1, restore_best_weights=True)
    reduce = ReduceLROnPlateau(monitor='val_loss', patience=3)
    csvlogger = CSVLogger(f"/kaggle/working/histories/history_model_{version}.csv")
    history = model.fit(train_data,
                         batch_size=batch_size,
                         epochs=epochs,
                         callbacks=[early,
                                    reduce,
                                    checkpoint,
                                    csvlogger,
                                    DisplayCallback()],
                         validation_data=val_data,
                         steps_per_epoch=len(df_train)//batch_size,
                         validation_steps=len(df_val)//batch_size,
                         validation_batch_size=batch_size
                        )
    return history
```

```
In [ ]: history_base_both = train_model_both(model_base_both, train_data_both, val_data_sg)
```

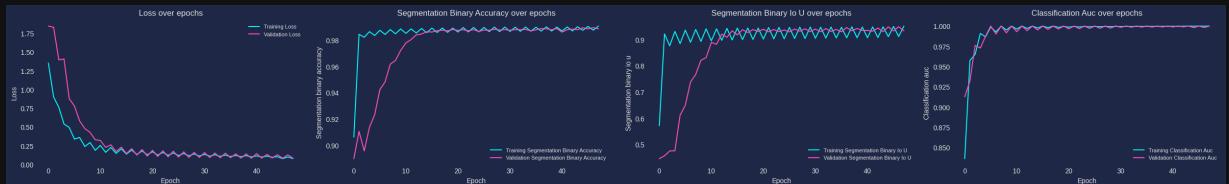
1/1 ————— 0s 26ms/step



Sample Prediction after epoch 48

```
70/70 ━━━━━━━━━━━━━━━━━━━━━━━━ 1s 10ms/step - classification_auc: 1.0000 - loss: 0.0758 - segmentation_binary_accuracy: 0.9907 - segmentation_binary_io_u: 0.9509 - val_classification_auc: 1.0000 - val_loss: 0.0811 - val_segmentation_binary_accuracy: 0.9886 - val_segmentation_binary_io_u: 0.9326 - learning_rate: 1.0000e-05
Epoch 48: early stopping
Restoring model weights from the end of the best epoch: 40.
```

```
In [ ]: show_history(version="seg_clf", parameters=[["loss", "val_loss"], ["segmentation_binary_accuracy"], ["segmentation_binary_io_u", "val_segmentation_binary_io_u"], ["classification_auc", "val_classification_auc"]])
```



Evaluating the baseline model for image segmentation and classification

```
In [ ]: model_base_both.evaluate(test_data_both, return_dict=True)
```

```
16/16 ━━━━━━━━━━━━━━━━━━━━━━━━ 2s 137ms/step - classification_auc: 0.9993 - loss: 0.1381 - segmentation_binary_accuracy: 0.9867 - segmentation_binary_io_u: 0.9380
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()` function when building your dataset.
    self.gen.throw(typ, value, traceback)
```

```
Out [ ]: {'classification_auc': 0.9992965459823608, 'loss': 0.13611197471618652, 'segmentation_binary_accuracy': 0.9866564869880676, 'segmentation_binary_io_u': 0.935377299785614}
```

Applying data augmentation technique

```
In [ ]: def data_transformation_both(image, mask, class_):
    img = random_brightness(image, max_delta=0.01, seed=seed_value)
    img = random_contrast(img, lower=0.5, upper=1, seed=seed_value)
    img = random_saturation(img, lower=0.5, upper=1, seed=seed_value)
    aux = tf.concat([img, mask], -1)

    aux_ = random_flip_left_right(aux, seed=seed_value)
    aux_ = random_flip_up_down(aux_, seed=seed_value)

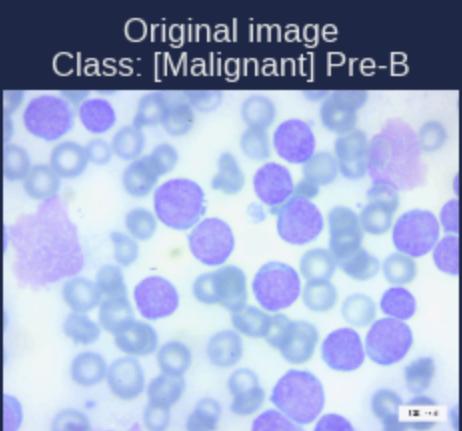
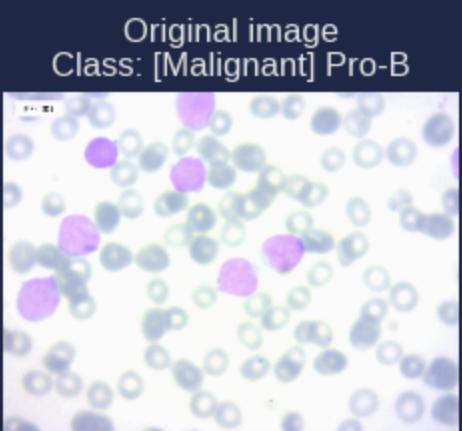
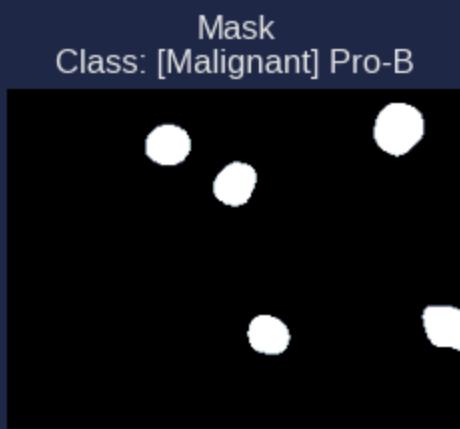
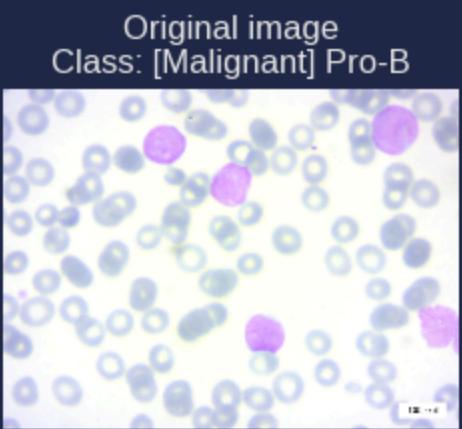
    class_ohe = tf.one_hot(indices=tf.squeeze(class_), axis=0), depth=4)

    return tf.cast(aux_[:, :, :3], tf.float32)/255, {"segmentation":tf.cast(aux_[:, :, 3], tf.int32)}, class_ohe

train_data_both_aug = train_data.map(data_transformation_both, num_parallel_calls=4)
train_data_both_aug = train_data_both_aug.batch(batch_size)
train_data_both_aug = train_data_both_aug.prefetch(autotune)
```

```
In [ ]: plt.subplots(4, 2, figsize=(5, 10))
for i, (image, pred) in enumerate(train_data_both_aug.take(1).unbatch()):
    mask, class_ = pred["segmentation"], pred["classification"]
    if i == 4:
        break
    plt.subplot(4, 2, (i*2)+1)
    plt.imshow(image)
    plt.axis("off")
    plt.title(f"Original image\nClass: {classes[np.argmax(class_)]}")

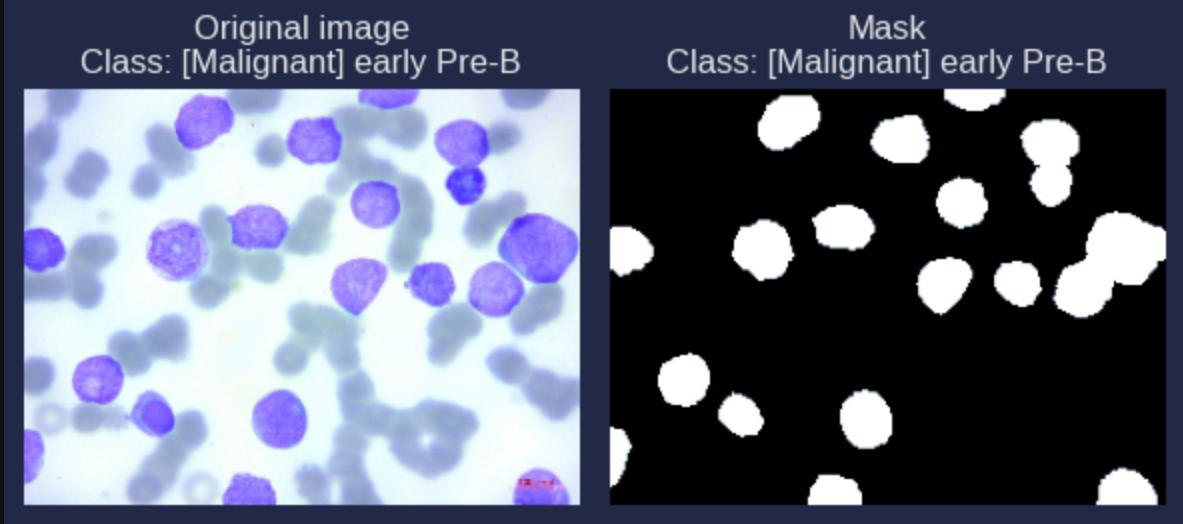
    plt.subplot(4, 2, (i*2)+2)
    plt.imshow(mask, cmap="bone")
    plt.axis("off")
    plt.title(f"Mask\nClass: {classes[np.argmax(class_)]}")
plt.tight_layout()
```



Let's train

```
In [ ]: model_base_both_aug = get_model_both()  
history_base_both_aug = train_model_both(model_base_both_aug, train_data_both)
```

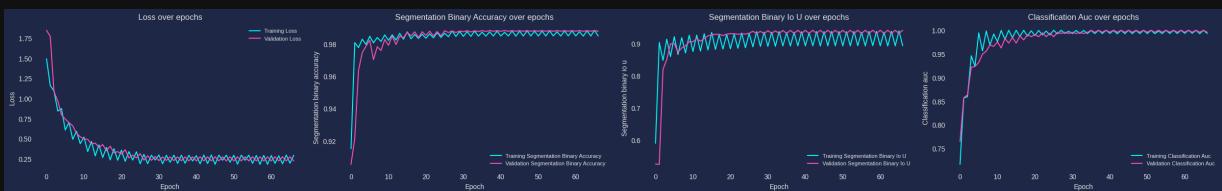
1/1 ————— 0s 27ms/step



Sample Prediction after epoch 67

```
70/70 ————— 24s 338ms/step - classification_auc:  
0.9935 - loss: 0.2861 - segmentation_binary_accuracy: 0.9849 - segmentation_b  
inary_io_u: 0.9139 - val_classification_auc: 0.9959 - val_loss: 0.2210 - val_s  
egmentation_binary_accuracy: 0.9877 - val_segmentation_binary_io_u: 0.9405 -  
learning_rate: 1.0000e-13  
Epoch 67: early stopping  
Restoring model weights from the end of the best epoch: 59.
```

```
In [ ]: show_history(version="seg_clf_aug", parameters=[["loss", "val_loss"],  
["segmentation_binary_accuracy",  
["segmentation_binary_io_u", "va  
["classification_auc", "val_class
```



Evaluating the baseline model for image segmentation and classification with DA

```
In [ ]: model_base_both_aug.evaluate(test_data_both, return_dict=True)
```

16/16 ————— 2s 137ms/step - classification_auc:
0.9916 - loss: 0.2695 - segmentation_binary_accuracy: 0.9858 - segmentation_b
inary_io_u: 0.9340

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of dat
a; interrupting training. Make sure that your dataset or generator can genera
te at least `steps_per_epoch * epochs` batches. You may need to use the `.`rep
eat()` function when building your dataset.
    self.gen.throw(typ, value, traceback)
```

```
Out[ ]: {'classification_auc': 0.9911665916442871,
         'loss': 0.2748297154903412,
         'segmentation_binary_accuracy': 0.9857146739959717,
         'segmentation_binary_io_u': 0.9311850666999817}
```

Evaluation metrics for both tasks

```
In [ ]: for i, (img, output) in enumerate(test_data_both):
            mask, class_ = output["segmentation"], output["classification"]
            aux = tf.reshape(mask, [-1])
            if i == 0:
                y_true_class = class_.numpy()
                y_true_mask = aux.numpy()
            else:
                y_true_class = np.concatenate((y_true_class, class_.numpy()), axis=0)
                y_true_mask = np.concatenate((y_true_mask, aux.numpy()), axis=0)

            scores_mask = []
            scores_class = []
            for model in [model_base_both, model_base_both_aug]:
                y_test_pred_proba = model.predict(test_data_both, verbose=0)

                y_test_pred_proba_class = y_test_pred_proba["classification"]
                y_test_pred_class = tf.math.round(y_test_pred_proba_class)
                scores_class.append(get_scores(y_true_class, y_test_pred_class, y_test_p

                y_test_pred_proba_mask = y_test_pred_proba["segmentation"]
                y_test_pred_proba_mask = tf.reshape(y_test_pred_proba_mask, [-1])
                y_test_pred_mask = tf.math.round(y_test_pred_proba_mask)
                score = get_scores(y_true_mask, y_test_pred_mask, y_test_pred_proba_mask)
                result_test = model.evaluate(test_data_both, return_dict=True)
                score["BinaryIOU"] = result_test["segmentation_binary_io_u"]
                score["BinaryAccuracy"] = result_test["segmentation_binary_accuracy"]
                scores_mask.append(score)
```

```
16/16 ----- 2s 134ms/step - classification_auc:
0.9993 - loss: 0.1381 - segmentation_binary_accuracy: 0.9867 - segmentation_b
inary_io_u: 0.9380
```

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of dat
a; interrupting training. Make sure that your dataset or generator can genera
te at least `steps_per_epoch * epochs` batches. You may need to use the `.`rep
eat()` function when building your dataset.
    self.gen.throw(typ, value, traceback)
```

```
16/16 ----- 2s 134ms/step - classification_auc:
0.9916 - loss: 0.2695 - segmentation_binary_accuracy: 0.9858 - segmentation_b
inary_io_u: 0.9340
```

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `repeat()` function when building your dataset.
  self.gen.throw(typ, value, traceback)
```

For classification task

```
In [ ]: comparison = pd.DataFrame(data=scores_class, index=["Baseline Model", "Baseline Model with DA"])
comparison.style.highlight_max(color = 'green', axis = 0).highlight_min(colo
```

	AUC_ROC	F1_Score	Accuracy	Precision	Recall
Baseline Model	0.999364	0.979420	0.977413	0.981942	0.977413
Baseline Model with DA	0.989956	0.932592	0.926078	0.943388	0.926078

For segmentation task

```
In [ ]: comparison = pd.DataFrame(data=scores_mask, index=["Baseline Model", "Baseline Model with DA"])
comparison.style.highlight_max(color = 'green', axis = 0).highlight_min(colo
```

	AUC_ROC	F1_Score	Accuracy	Precision	Recall	BinaryIOU	BinaryAccuracy
Baseline Model	0.988337	0.986632	0.986656	0.986612	0.986656	0.935377	0.986656
Baseline Model with DA	0.987248	0.985703	0.985715	0.985692	0.985715	0.931185	0.985715

Confusion Matrix for classification

```
In [ ]: for i, (img, output) in enumerate(test_data_both):
    class_ = tf.argmax(output["classification"], axis=1)
    if i == 0:
        y_true = class_.numpy()
    else:
        y_true = np.concatenate((y_true, class_.numpy()), axis=0)

y_test_pred_proba = model_base_both.predict(test_data_both, verbose=0)
y_test_pred_proba = y_test_pred_proba["classification"]
y_test_pred = tf.argmax(y_test_pred_proba, axis=1)

def my_cm(y_true, y_pred, title):
    cm_val = confusion_matrix(y_true, y_pred)
    cm_pgs = np.round(confusion_matrix(y_true, y_pred, normalize='true')*100

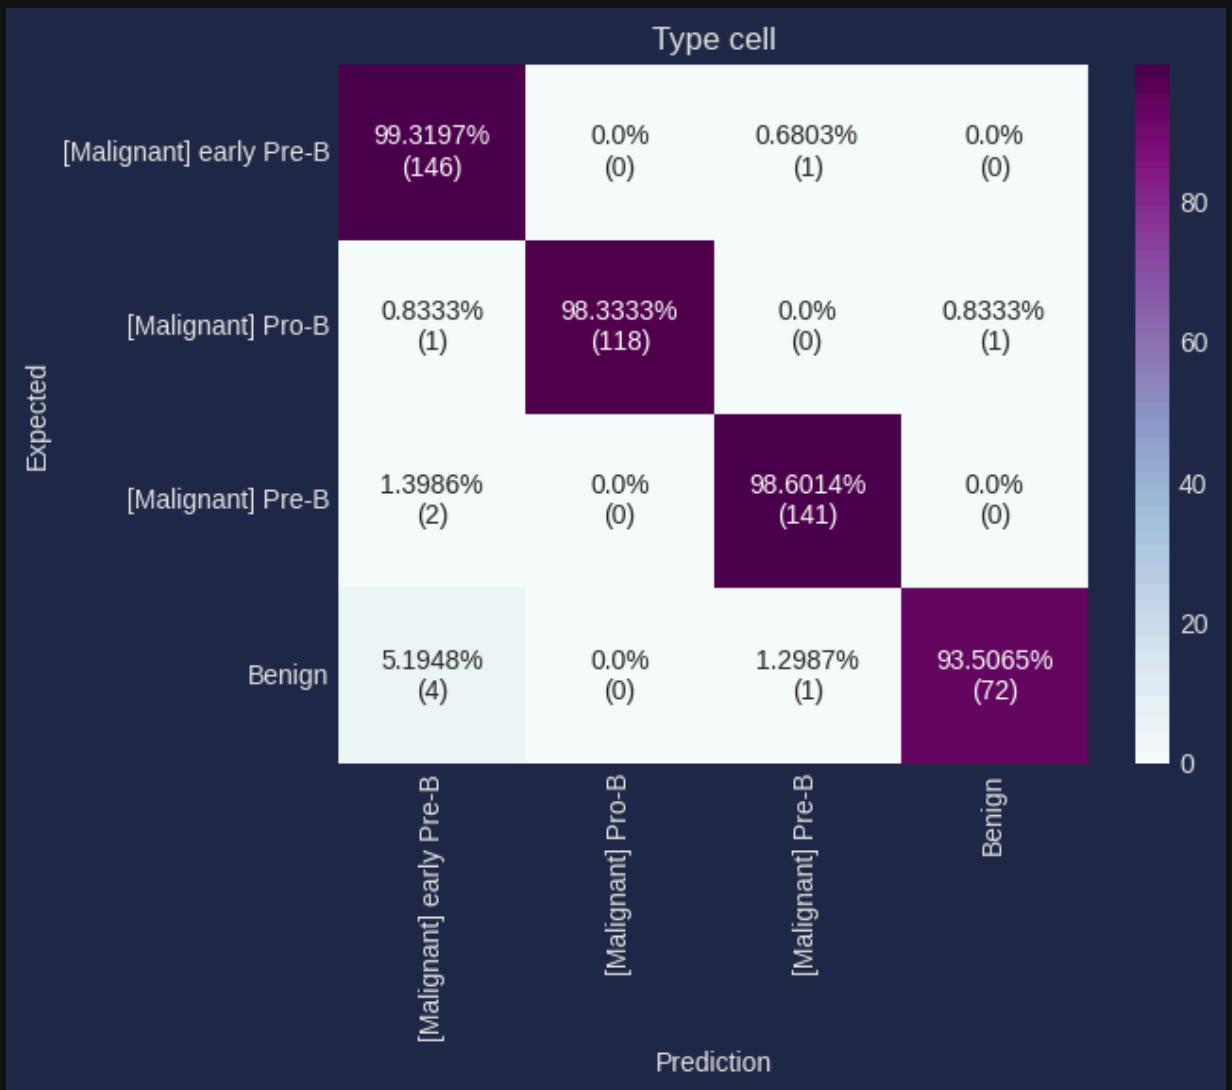
    formatted_text = (np.asarray([f'{pgs}%\n{val}' for val, pgs in zip(cm_
```

```

        plt.xlabel("Prediction")
        plt.ylabel("Expected")
        return

my_cm(y_true, y_test_pred, title="Type cell")

```



Let's take a look at the predictions.

```

In [ ]: plt.subplots(4, 4, figsize=(15, 13))
for i, (image, output) in enumerate(test_data_both.shuffle(100).take(1).unbatch()):
    if i == 4:
        break

    pred = model_base_both.predict(tf.expand_dims(image, axis=0), verbose=0)

    class_ = np.argmax(output["classification"])
    pred_class_proba = pred["classification"]
    pred_class = tf.argmax(pred_class_proba, axis=1)

    mask = output["segmentation"]
    pred_mask = tf.squeeze(pred["segmentation"], axis=0)

```

```

image_cropped = apply_mask(tf.cast(image*255, tf.uint8).numpy(), tf.cast

plt.subplot(4, 4, (i*4)+1)
plt.imshow(image)
plt.axis("off")
plt.title(f"Original Image\nExpected Class: {classes[class_]}")

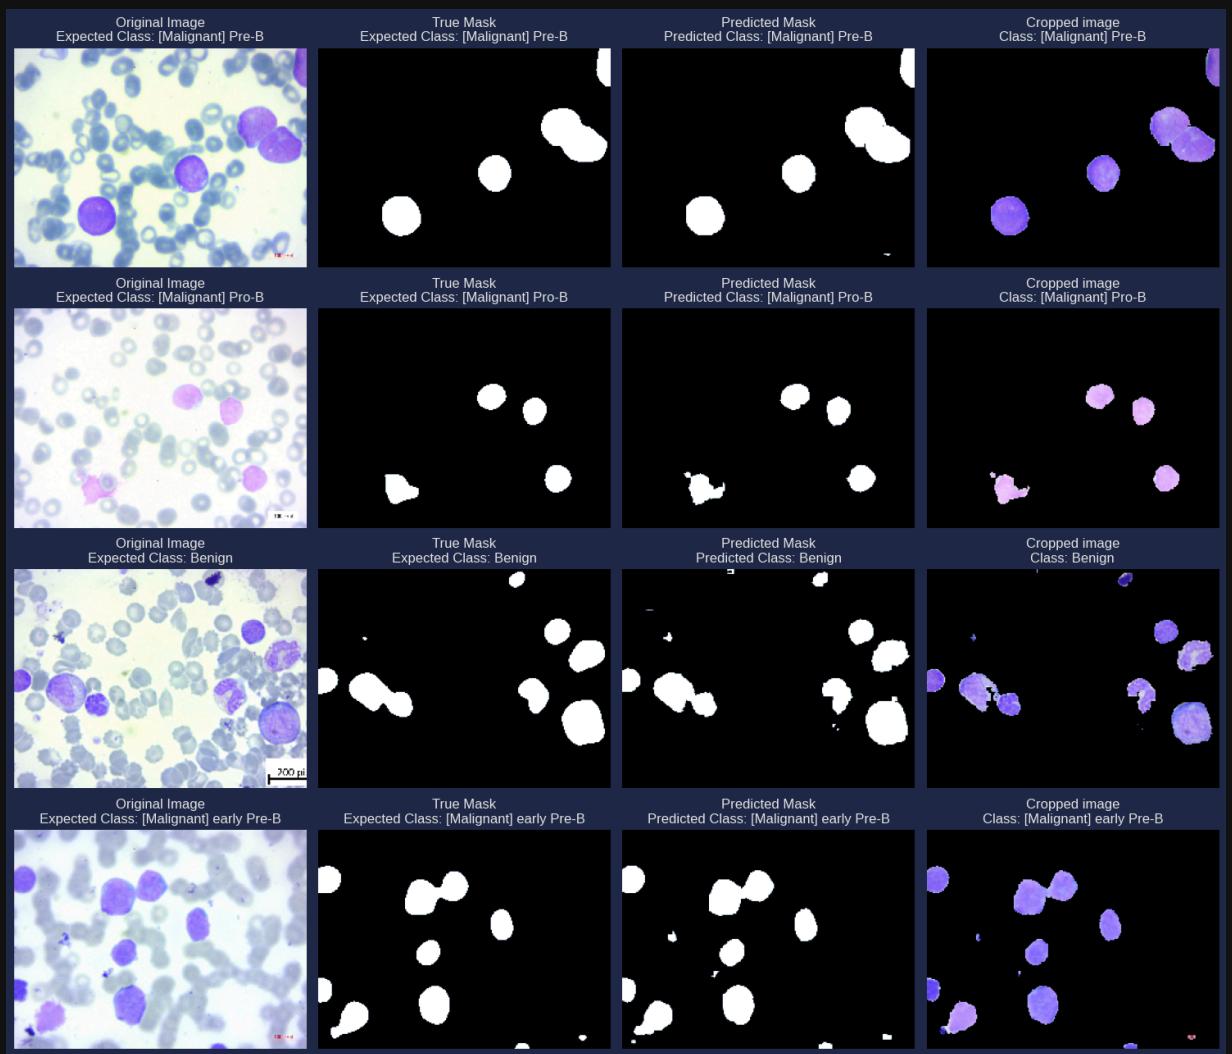
plt.subplot(4, 4, (i*4)+2)
plt.imshow(mask, cmap="bone")
plt.axis("off")
plt.title(f"True Mask\nExpected Class: {classes[class_]}")

plt.subplot(4, 4, (i*4)+3)
plt.imshow(pred_mask, cmap="bone")
plt.axis("off")
plt.title(f"Predicted Mask\nPredicted Class: {classes[int(pred_class)]}")

plt.subplot(4, 4, (i*4)+4)
plt.imshow(image_cropped)
plt.axis("off")
plt.title(f"Cropped image\nClass: {classes[int(pred_class)]}")

plt.tight_layout()

```



Converting Notebook to PDF

```
In [ ]: !apt-get install texlive-xetex texlive-fonts-recommended texlive-plain-gener  
!pip install pypandoc nbconvert [webpdf]  
!playwright install  
!playwright install-deps  
  
from google.colab import drive  
from IPython.display import clear_output  
drive.mount('/content/drive')  
clear_output(wait=False)  
  
In [ ]: %%capture  
!jupyter nbconvert --to webpdf /content/drive/MyDrive/ML_Zoomcamps_2024/Caps
```

Thank you for your time! 😊