


Estimation of obesity levels based on eating habits and physical condition - [DataTalks.Club](#)'s midterm project by Alexander D. Rios

 [Open in Colab](#)



Obesity is a global problem, and an individual's eating habits and physical condition can provide insights into their health status. Poor eating habits and a lack of physical activity can lead to severe obesity, which may have fatal consequences.

In this project, I used data on individuals from Colombia, Peru, and Mexico obtained from the [UC Irvine Machine Learning Repository](#). The dataset contains 17 attributes and 2,111 records, with each record labeled using the class variable `NObesity` (Obesity Level). This variable categorizes obesity levels into the following classes:

- Insufficient Weight
- Normal Weight
- Overweight Level I
- Overweight Level II
- Obesity Type I
- Obesity Type II
- Obesity Type III

Seventy-seven percent of the data was generated synthetically using the Weka tool with the SMOTE filter, while 23% was collected directly from users via a web platform.

Downloading the dataset

```
In [5]: !wget https://archive.ics.uci.edu/static/public/544/estimation+of+obesity+le
!unzip -o /content/estimation+of+obesity+levels+based+on+eating+habits+and+p

--2024-11-25 14:16:53-- https://archive.ics.uci.edu/static/public/544/estima
tion+of+obesity+levels+based+on+eating+habits+and+physical+condition.zip
Resolving archive.ics.uci.edu (archive.ics.uci.edu)... 128.195.10.252
Connecting to archive.ics.uci.edu (archive.ics.uci.edu)|128.195.10.252|:44
3... connected.
HTTP request sent, awaiting response... 200 OK
Length: unspecified
Saving to: 'estimation+of+obesity+levels+based+on+eating+habits+and+physical+
condition.zip'

estimation+of+obesi      [ <=>                ] 56.32K  --.-KB/s    in 0.1s

2024-11-25 14:16:53 (433 KB/s) - 'estimation+of+obesity+levels+based+on+eatin
g+habits+and+physical+condition.zip' saved [57676]

Archive:  /content/estimation+of+obesity+levels+based+on+eating+habits+and+ph
ysical+condition.zip
  inflating: ObesityDataSet_raw_and_data_sinthetic.csv
```

Installing some packages

```
In [56]: !pip install mplcyberpunk
!pip install catboost

from IPython.display import clear_output

clear_output(wait=False)
```

Feature description

Feature	Description
Gender	
Age	
Height	
Weight	
family_history_with_overweight	Has a family member suffered or suffers from overweight?
FAVC	Do you eat high caloric food frequently?
FCVC	Do you usually eat vegetables in your meals?
NCP	How many main meals do you have daily?
CAEC	Do you eat any food between meals?

SMOKE	Do you smoke?
CH20	How much water do you drink daily?
SCC	Do you monitor the calories you eat daily?
FAF	How often do you have physical activity?
TUE	How much time do you use technological devices such as cell phone, videogames, television, computer and others?
CALC	How often do you drink alcohol?
MTRANS	Which transportation do you usually use?
NObeyesdad	Obesity level

Loading the dataset

```
In [7]: import pandas as pd
df_raw = pd.read_csv("/content/ObesityDataSet_raw_and_data_sinthetic.csv")
df_raw.head()
```

```
Out [7]:
```

	Gender	Age	Height	Weight	family_history_with_overweight	FAVC	FCVC	NCP	
0	Female	21.0	1.62	64.0	yes	no	2.0	3.0	Son
1	Female	21.0	1.52	56.0	yes	no	3.0	3.0	Son
2	Male	23.0	1.80	77.0	yes	no	2.0	3.0	Son
3	Male	27.0	1.80	87.0	no	no	3.0	3.0	Son
4	Male	22.0	1.78	89.8	no	no	2.0	1.0	Son

Exploratory data analysis (EDA)

Renaming the columns

```
In [8]: df_raw = df_raw.rename(columns={"family_history_with_overweight": "overweigh",
                                     "FAVC": "eat_HC_food",
                                     "FCVC": "eat_vegetables",
                                     "NCP": "main_meals",
                                     "CAEC": "snack",
                                     "CH20": "drink_water",
                                     "SCC": "monitoring_calories",
                                     "FAF": "physical_activity",
                                     "TUE": "use_of_technology",
                                     "CALC": "drink_alcohol",
                                     "MTRANS": "transportation_type",
                                     "NObeyesdad": "obesity_level"
                                     })
```

```
}).rename(columns=str.lower)  
df_raw.columns
```

```
Out[8]: Index(['gender', 'age', 'height', 'weight', 'overweight_familiar',  
              'eat_hc_food', 'eat_vegetables', 'main_meals', 'snack', 'smoke',  
              'drink_water', 'monitoring_calories', 'physical_activity',  
              'use_of_technology', 'drink_alcohol', 'transportation_type',  
              'obesity_level'],  
           dtype='object')
```

Dropping duplicated

```
In [9]: df_raw.duplicated().sum()
```

```
Out[9]: 24
```

```
In [10]: df_raw = df_raw.drop_duplicates()  
df_raw.duplicated().sum()
```

```
Out[10]: 0
```

Looking for missing values

```
In [11]: df_raw.isna().sum()
```

Out [11]:

	0
gender	0
age	0
height	0
weight	0
overweight_familiar	0
eat_hc_food	0
eat_vegetables	0
main_meals	0
snack	0
smoke	0
drink_water	0
monitoring_calories	0
physical_activity	0
use_of_technology	0
drink_alcohol	0
transportation_type	0
obesity_level	0

dtype: int64

Statistical description

```
In [12]: df_raw.describe(include="all").T
```

```
Out [12]:
```

	count	unique	top	freq	mean	std	min
gender	2087	2	Male	1052	NaN	NaN	NaN
age	2087.0	NaN	NaN	NaN	24.35309	6.368801	14.0
height	2087.0	NaN	NaN	NaN	1.702674	0.093186	1.45
weight	2087.0	NaN	NaN	NaN	86.85873	26.190847	39.0
overweight_familiar	2087	2	yes	1722	NaN	NaN	NaN
eat_hc_food	2087	2	yes	1844	NaN	NaN	NaN
eat_vegetables	2087.0	NaN	NaN	NaN	2.421466	0.534737	1.0
main_meals	2087.0	NaN	NaN	NaN	2.701179	0.764614	1.0
snack	2087	4	Sometimes	1761	NaN	NaN	NaN
smoke	2087	2	no	2043	NaN	NaN	NaN
drink_water	2087.0	NaN	NaN	NaN	2.004749	0.608284	1.0
monitoring_calories	2087	2	no	1991	NaN	NaN	NaN
physical_activity	2087.0	NaN	NaN	NaN	1.012812	0.853475	0.0
use_of_technology	2087.0	NaN	NaN	NaN	0.663035	0.608153	0.0
drink_alcohol	2087	4	Sometimes	1380	NaN	NaN	NaN
transportation_type	2087	5	Public_Transportation	1558	NaN	NaN	NaN
obesity_level	2087	7	Obesity_Type_I	351	NaN	NaN	NaN

Looking for the classes in categorical features

```
In [13]: df_raw.head()
```

```
Out [13]:
```

	gender	age	height	weight	overweight_familiar	eat_hc_food	eat_vegetables	main_m
0	Female	21.0	1.62	64.0	yes	no	2.0	
1	Female	21.0	1.52	56.0	yes	no	3.0	
2	Male	23.0	1.80	77.0	yes	no	2.0	
3	Male	27.0	1.80	87.0	no	no	3.0	
4	Male	22.0	1.78	89.8	no	no	2.0	

```
In [14]: df_raw.select_dtypes("object").nunique()
```

Out [14]: 0

gender	2
overweight_familiar	2
eat_hc_food	2
snack	4
smoke	2
monitoring_calories	2
drink_alcohol	4
transportation_type	5
obesity_level	7

dtype: int64

```
In [15]: df = df_raw.copy()
```


Data visualization

```
In [16]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import mplcyberpunk

plt.style.use("cyberpunk")
```


Balance of the target variable

```
In [17]: plt.figure(figsize=(15, 5))
sns.histplot(df, y="obesity_level", alpha=1)
plt.ylabel("Obesity level");
```



The relationship between obesity levels and the gender

```
In [18]: plt.figure(figsize=(15, 5))
sns.histplot(df, y="obesity_level", hue="gender", multiple="dodge", shrink=0
plt.ylabel("Obesity level");
```



The balance in the gender feature:

```
In [19]: df.gender.value_counts(normalize=True)
```

Out [19]:

	proportion
--	------------

gender	
--------	--

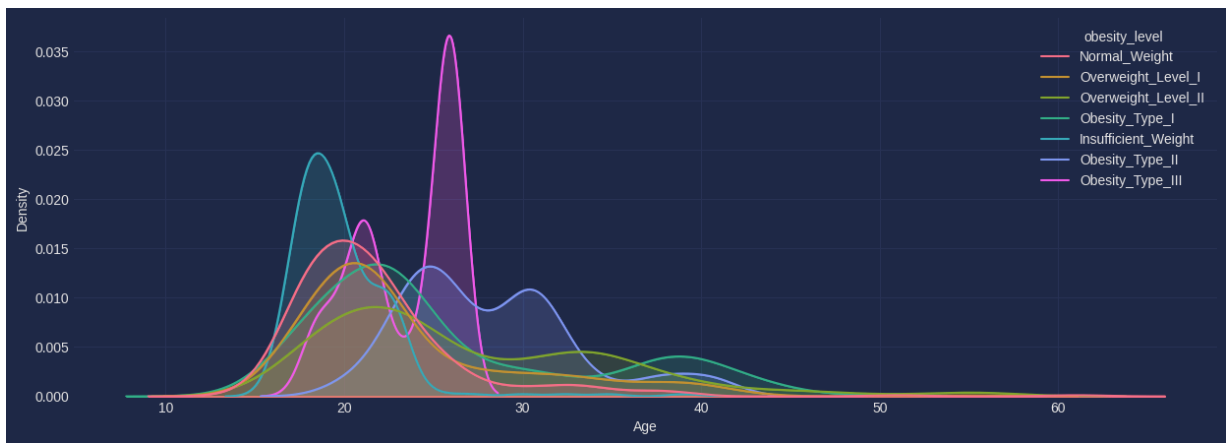
Male	0.504073
------	----------

Female	0.495927
--------	----------

dtype: float64

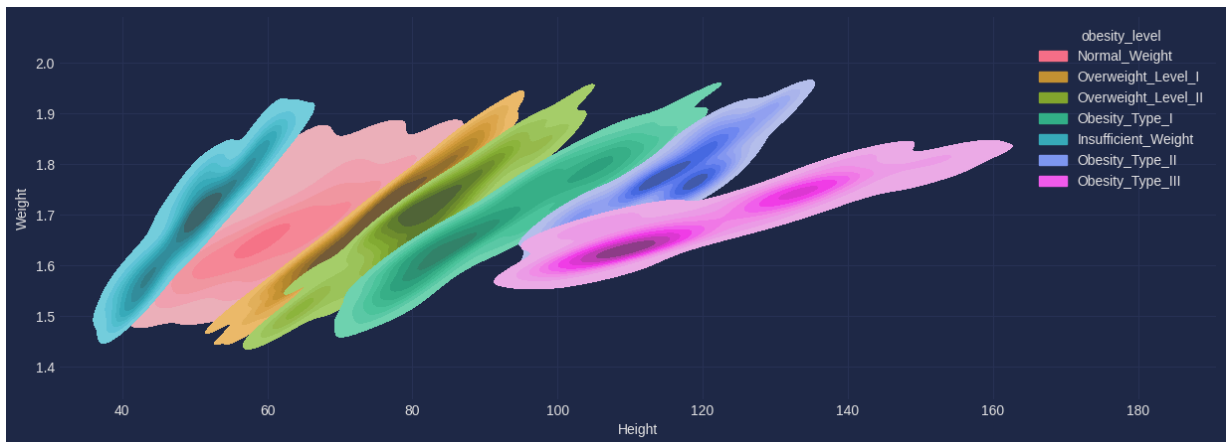
The distribution of age in relation to obesity levels

```
In [20]: plt.figure(figsize=(15, 5))
sns.kdeplot(df, x="age", hue="obesity_level", fill=True, alpha=.2)
sns.kdeplot(df, x="age", hue="obesity_level")
plt.xlabel("Age");
```


Height and weight in relation to obesity levels

```
In [21]: plt.figure(figsize=(15, 5))
sns.kdeplot(data=df, x="weight", y="height", hue="obesity_level", fill=True)
plt.xlabel("Height")
plt.ylabel("Weight");
```




The relationship between obesity levels and having overweight family members

```
In [22]: plt.figure(figsize=(15, 5))
sns.histplot(df, y="obesity_level", hue="overweight_familiar", multiple="dod")
plt.ylabel("Obesity level");
```




The relationship between obesity levels and the consumption of high-calorie foods

```
In [23]: plt.figure(figsize=(15, 5))
sns.histplot(df, y="obesity_level", hue="eat_hc_food", multiple="dodge", shrink=0.5)
plt.ylabel("Obesity level");
```




The relationship between obesity levels and vegetable eating habits

```
In [24]: plt.figure(figsize=(15, 5))
sns.violinplot(data=df, x="obesity_level", y="eat_vegetables", hue="obesity_level")
plt.ylabel("Vegetable eating habits")
plt.xlabel("Obesity level");
```




The relationship between obesity levels and the number of main meals

```
In [25]: plt.figure(figsize=(15, 5))
sns.violinplot(data=df, x="obesity_level", y="main_meals", hue="obesity_level")
plt.ylabel("Main meals")
plt.xlabel("Obesity level");
```




The relationship between obesity levels and snack eating habits

```
In [26]: plt.figure(figsize=(15, 8))
sns.histplot(df, y="obesity_level", hue="snack", multiple="dodge", shrink=0.
plt.ylabel("Obesity level");
```



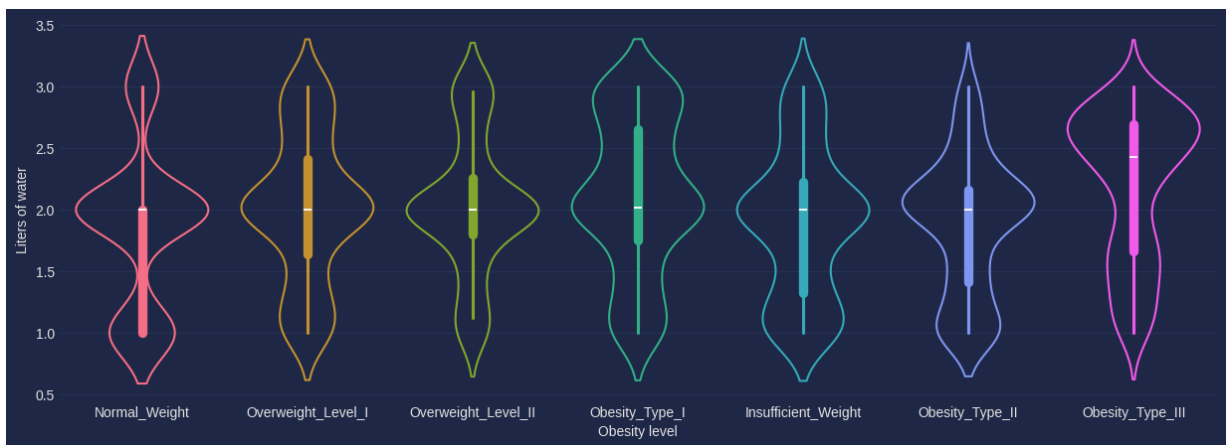
The relationship between obesity levels and smoking habits

```
In [27]: plt.figure(figsize=(15, 5))
sns.histplot(df, y="obesity_level", hue="smoke", multiple="dodge", shrink=0.
plt.ylabel("Obesity level");
```




The relationship between obesity levels and water consumption

```
In [28]: plt.figure(figsize=(15, 5))
sns.violinplot(data=df, x="obesity_level", y="drink_water", hue="obesity_level")
plt.ylabel("Liters of water")
plt.xlabel("Obesity level");
```




Calorie monitoring habits in relation to obesity levels

```
In [29]: plt.figure(figsize=(15, 5))
sns.histplot(df, y="obesity_level", hue="monitoring_calories", multiple="dod")
plt.ylabel("Obesity level");
```



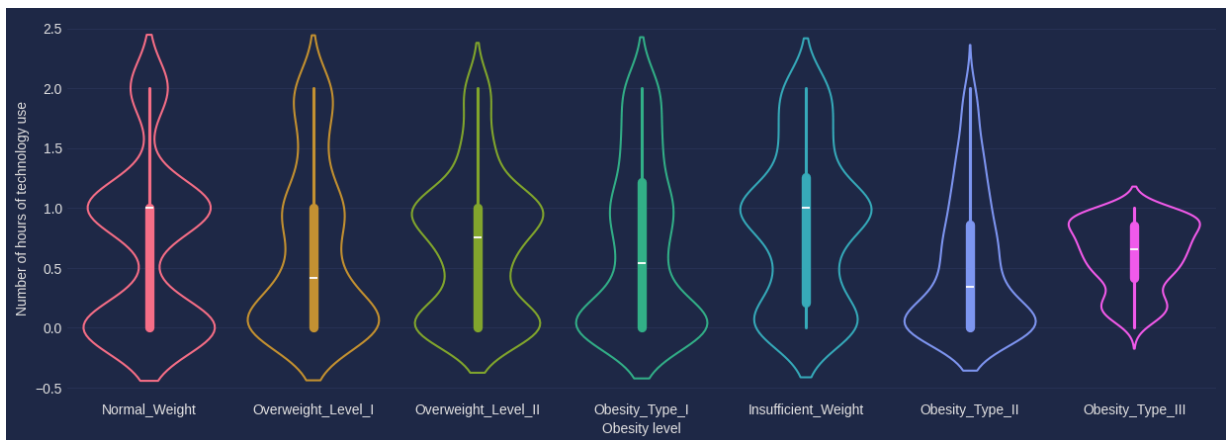
Physical activity in relation to obesity levels

```
In [30]: plt.figure(figsize=(15, 5))
sns.violinplot(data=df, x="obesity_level", y="physical_activity", hue="obesi
plt.ylabel("Number of days of physical activity")
plt.xlabel("Obesity level");
```




Technology usage habits in relation to obesity levels

```
In [31]: plt.figure(figsize=(15, 5))
sns.violinplot(data=df, x="obesity_level", y="use_of_technology", hue="obesi
plt.ylabel("Number of hours of technology use")
plt.xlabel("Obesity level");
```




Alcohol drinking habits in relation to obesity levels

```
In [32]: plt.figure(figsize=(15, 8))
sns.histplot(df, y="obesity_level", hue="drink_alcohol", multiple="dodge", s
plt.xlabel("Obesity level");
```



Obesity levels according to the type of transportation used

```
In [33]: plt.figure(figsize=(15, 8))
sns.histplot(df, y="obesity_level", hue="transportation_type", multiple="dod
plt.xlabel("Obesity level");
```



Splitting the dataset

Setting seeds for reproducibility

```
In [34]: import random
import os
import keras
```

```
seed_value = 42
os.environ['PYTHONHASHSEED'] = str(seed_value)
random.seed(seed_value)
np.random.seed(seed_value)
```

Setting the validation framework

```
In [35]: from sklearn.model_selection import train_test_split

df_full_train, df_test = train_test_split(df, test_size=0.15, random_state=s
df_train, df_val = train_test_split(df_full_train, test_size=0.15, random_st
len(df_train), len(df_test), len(df_val))
```

```
Out [35]: (1507, 314, 266)
```

Removing specific columns and separating features from the target column

According to [the dataset's article](#), the data was labeled using the following equation:

$$\text{Mass body index} = \frac{\text{weight}}{\text{height}^2}$$

Therefore, we need to delete at least a feature into previous equation.

```
In [36]: cols_drop = ["obesity_level", "weight"]
X_full_train, y_full_train = df_full_train.drop(cols_drop, axis=1), df_full_
X_train, y_train = df_train.drop(cols_drop, axis=1), df_train["obesity_leve
X_val, y_val = df_val.drop(cols_drop, axis=1), df_val["obesity_level"]
X_test, y_test = df_test.drop(cols_drop, axis=1), df_test["obesity_level"]
```

Standardization

```
In [37]: from sklearn.preprocessing import StandardScaler

numeric_cols = X_train.select_dtypes(exclude=["object"]).columns

ss = StandardScaler().set_output(transform="pandas")

ss.fit(X_train[numeric_cols])
ss.transform(X_train[numeric_cols])
```

Out [37]:

	age	height	eat_vegetables	main_meals	drink_water	physical_activity	u
2071	-0.859706	0.482798	1.091725	0.378982	0.667080	-0.024754	
190	-0.679409	-1.084919	-0.803418	0.378982	-0.005343	-0.007330	
1361	-0.996350	-0.153529	0.593585	0.378982	-0.009926	0.410323	
2109	0.011829	0.409630	1.091725	0.378982	1.387169	0.157412	
325	-0.520938	-1.620791	-0.803418	-2.235837	-1.639098	-0.007330	
...	
1357	-0.996350	0.891401	-0.803418	-0.306828	1.066065	-0.007330	
1013	4.906080	0.729214	-0.803418	0.378982	-0.005343	-0.007330	
1971	-0.790813	1.243673	1.091725	0.378982	1.137754	0.609467	
1265	-1.048027	0.399074	-0.803418	0.378982	-0.005343	-1.191614	
1881	0.029785	-0.069697	1.091725	0.378982	1.177656	-0.792753	

1507 rows × 7 columns

One-hot encoding of features

In [38]:

```

from sklearn.feature_extraction import DictVectorizer

dict_X_full_train = X_full_train.to_dict("records")
dict_X_train = X_train.to_dict("records")
dict_X_val = X_val.to_dict("records")
dict_X_test = X_test.to_dict("records")

dv = DictVectorizer(sparse=False).set_output(transform="pandas")
dv.fit(dict_X_train)

dv.transform(dict_X_train).head()

```

Out [38]:

	age	drink_alcohol=Always	drink_alcohol=Frequently	drink_alcohol=Sometimes	u
0	18.862264	0.0	0.0	1.0	
1	20.000000	0.0	0.0	1.0	
2	18.000000	0.0	0.0	1.0	
3	24.361936	0.0	0.0	1.0	
4	21.000000	0.0	0.0	1.0	

5 rows × 5 columns

Creating a Pipeline

```
In [39]: from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.pipeline import Pipeline

class MyStandardScaler(BaseEstimator, TransformerMixin):
    def __init__(self, numeric_cols):
        self.ss = StandardScaler().set_output(transform="pandas")
        self.numeric_cols = numeric_cols
        return

    def fit(self, X):
        self.ss.fit(X[self.numeric_cols])
        return self

    def transform(self, X):
        X[self.numeric_cols] = self.ss.transform(X[self.numeric_cols])
        return X.to_dict("records")

numeric_cols = X_train.select_dtypes(exclude=["object"]).columns
pipe = Pipeline([('ss', MyStandardScaler(numeric_cols=numeric_cols)), ('dv',

X_train = pipe.fit_transform(X_train)
X_full_train = pipe.transform(X_full_train)
X_val = pipe.transform(X_val)
X_test = pipe.transform(X_test)
```

Label encoding

```
In [40]: from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
le.fit(y_train)
y_full_train = le.transform(y_full_train)
y_train = le.transform(y_train)
y_val = le.transform(y_val)
y_test = le.transform(y_test)
```

Computing weights for classes and samples

```
In [41]: from sklearn.utils.class_weight import compute_sample_weight, compute_class_

class_full_weight = compute_class_weight(class_weight="balanced", classes=np
class_full_weight = dict(zip(np.unique(y_full_train), class_full_weight))

class_weight = compute_class_weight(class_weight="balanced", classes=np.uniq
class_weight = dict(zip(np.unique(y_train), class_weight))

sample_full_weights = compute_sample_weight(
```

```

        class_weight=class_full_weight,
        y=y_full_train
    )

    sample_weights = compute_sample_weight(
        class_weight=class_weight,
        y=y_train
    )

```

Training the models

Logistic Regression

```
In [42]: from sklearn.metrics import roc_auc_score
```

```
In [ ]: from sklearn.linear_model import LogisticRegression
        from sklearn.model_selection import GridSearchCV

        parameters = {"max_iter":[200, 300, 400, 500],
                       "C":[20, 15, 10],
                       "class_weight":["balanced"],
                       "solver":["lbfgs", "newton-cg", "sag", "saga"]}

        lr = LogisticRegression(random_state=seed_value)
        gs_lr = GridSearchCV(lr, param_grid=parameters, n_jobs=-1, cv=5, scoring="ro
        gs_lr.fit(X_full_train, y_full_train, sample_weight=sample_full_weights)
```

```
/usr/local/lib/python3.10/dist-packages/numpy/ma/core.py:2820: RuntimeWarnin
g: invalid value encountered in cast
    _data = np.array(data, dtype=dtype, copy=copy,
```

```
Out [ ]:  ▸ GridSearchCV ⓘ ⓘ
          ▸ best_estimator_: LogisticRegression
            ▸ LogisticRegression ⓘ
```

```
In [ ]: y_pred_test_lr = gs_lr.predict_proba(X_test)
        roc_auc_score(y_test, y_pred_test_lr, multi_class="ovr")
```

```
Out [ ]: 0.889393203661662
```


Decision Tree

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
        parameters = {"criterion":["gini", "entropy", "log_loss"],
                       "max_depth":[5, 6, 7, 8, 9],
                       "max_leaf_nodes":[30, 40, 50, 60],
                       "min_samples_split":[30, 40],
```

```
"class_weight":["balanced"],}
```

```
dt = DecisionTreeClassifier(random_state=seed_value)
gs_dt = GridSearchCV(dt, param_grid=parameters, n_jobs=-1, cv=5, scoring="roc_auc")
gs_dt.fit(X_full_train, y_full_train, sample_weight=sample_full_weights)
```

Out []:



In []:

```
y_pred_test_dt = gs_dt.predict_proba(X_test)
roc_auc_score(y_test, y_pred_test_dt, multi_class="ovr")
```

Out []: 0.9213027545513881

Random Forest


In []:

```
from sklearn.ensemble import RandomForestClassifier
parameters = {"criterion":["gini", "entropy", "log_loss"],
              "n_estimators":[100, 200, 300, 400],
              "max_depth":[5, 6, 7, 8, 9],
              "class_weight":["balanced"],}

rf = RandomForestClassifier(random_state=seed_value)
gs_rf = GridSearchCV(rf, param_grid=parameters, n_jobs=-1, cv=5, scoring="roc_auc")
gs_rf.fit(X_full_train, y_full_train, sample_weight=sample_full_weights)
```

```
/usr/local/lib/python3.10/dist-packages/numpy/ma/core.py:2820: RuntimeWarning:
invalid value encountered in cast
  _data = np.array(data, dtype=dtype, copy=copy,
```

Out []:



In []:

```
y_pred_test_rf = gs_rf.predict_proba(X_test)
roc_auc_score(y_test, y_pred_test_rf, multi_class="ovr")
```

Out []: 0.9721096055019289

Gradient Boosting with XGBoost

In [55]:

```
import xgboost as xgb

dtrain = xgb.DMatrix(X_train, label=y_train, feature_names=X_train.columns.tolist())
```

```

dval = xgb.DMatrix(X_val, label=y_val, feature_names=X_val.columns.tolist())
dtest = xgb.DMatrix(X_test, label=y_test, feature_names=X_test.columns.tolist())

xgb_params = {
    'eta': 0.1,
    'max_depth': 6,
    'gamma': 0.0001,
    'min_child_weight': 1,
    'alpha': 0.01,

    'objective': 'multi:softprob',
    'num_class': 7,
    'nthread': 8,
    'eval_metric': 'auc',
    'num_parallel_tree': 5,

    'seed': seed_value,
    'verbosity': 1,
}

watchlist = [(dtrain, 'train'), (dval, 'val')]
xgb_clf = xgb.train(xgb_params, dtrain, num_boost_round=10000, early_stopping_rounds=10)
y_test_pred_xgb = xgb_clf.predict(dtest)
print(f"ROC: {roc_auc_score(y_test, y_test_pred_xgb, multi_class='ovr')}")

[0]      train-auc:0.96308      val-auc:0.88976
[200]     train-auc:1.00000      val-auc:0.97088
[400]     train-auc:1.00000      val-auc:0.97176
[600]     train-auc:1.00000      val-auc:0.97269
[800]     train-auc:1.00000      val-auc:0.97309
[942]     train-auc:1.00000      val-auc:0.97303
ROC:0.9806772041514711

```

Gradient Boosting with CatBoost

According to [this blog](#).

CatBoost operates on the principle of gradient boosting, where it builds the model in a stage-wise fashion. It starts with a simple model and incrementally improves it by adding new models that correct the errors made by the preceding ones.

CatBoost introduces several key innovations:

Ordered Boosting

One of the core innovations of CatBoost is its ordered boosting mechanism. Traditional gradient boosting methods can suffer from prediction shift due to the overlap between the training data for the base models and the data used to calculate the gradients. CatBoost addresses this by introducing a random permutation of the dataset in each iteration and using only the data before each example in the permutation for training. This approach reduces overfitting and improves model robustness.

Symmetric Trees

CatBoost builds balanced trees, also known as symmetric trees, as its base predictors. Unlike traditional gradient boosting methods that build trees leaf-wise or depth-wise, CatBoost's symmetric trees ensure that all leaf nodes at the same level share the same decision rule. This leads to faster execution and reduces the likelihood of overfitting.

```
In [53]: from catboost import CatBoostClassifier

cbc = CatBoostClassifier(loss_function='MultiClass',
                        eval_metric='AUC',
                        iterations=5000,
                        depth=6,
                        classes_count=7,
                        class_weights=class_weight,
                        learning_rate=0.1,
                        od_type='Iter',
                        early_stopping_rounds=1000,
                        bootstrap_type='MVS',
                        sampling_frequency='PerTree',
                        random_seed=seed_value,
                        verbose=200)

cbc.fit(X_train, y_train, sample_weight=sample_weights, eval_set=(X_val, y_val),
        y_test_pred_cbc = cbc.predict_proba(X_test)
print(f"ROC:{roc_auc_score(y_test, y_test_pred_cbc, multi_class='ovr')}")
```

0:	test: 0.8701799	best: 0.8701799 (0)	total: 36.5ms	remaining: 6m
5s				
200:	test: 0.9843814	best: 0.9843814 (200)	total: 6.27s	remaining: 5m
5s				
400:	test: 0.9859230	best: 0.9860489 (384)	total: 9.02s	remaining: 3m
35s				
600:	test: 0.9859990	best: 0.9862174 (434)	total: 11.7s	remaining: 3m
2s				
800:	test: 0.9860843	best: 0.9865136 (706)	total: 14.4s	remaining: 2m
45s				
1000:	test: 0.9861372	best: 0.9865136 (706)	total: 19s	remaining: 2m
51s				
1200:	test: 0.9859280	best: 0.9865136 (706)	total: 21.9s	remaining: 2m
40s				
1400:	test: 0.9860076	best: 0.9865136 (706)	total: 24.5s	remaining: 2m
30s				
1600:	test: 0.9857638	best: 0.9865136 (706)	total: 27.2s	remaining: 2m
22s				

Stopped by overfitting detector (1000 iterations wait)

```
bestTest = 0.9865136358
bestIteration = 706

Shrink model to first 707 iterations.
ROC:0.9849353946292636
```

Neural Networks

```
In [ ]: import tensorflow as tf
        from tensorflow.data import Dataset
        from tensorflow.keras import Input, Model
        from tensorflow.keras.layers import Dense, Dropout
        from tensorflow.keras.callbacks import EarlyStopping
        from tensorflow.keras.losses import SparseCategoricalCrossentropy
        from tensorflow.keras.optimizers import Adamax
        from tensorflow.keras.metrics import SparseCategoricalAccuracy

In [ ]: tf.keras.backend.clear_session()
        keras.utils.set_random_seed(seed_value)

        autotune = tf.data.AUTOTUNE

        batch_size = 32

        train_data = Dataset.from_tensor_slices((X_train.astype(float), y_train.astype(float)))
        val_data = Dataset.from_tensor_slices((X_val.astype(float), y_val.astype(float)))
        test_data = Dataset.from_tensor_slices((X_test.astype(float), y_test.astype(float)))

        early = EarlyStopping(monitor='val_loss', patience=10)

        inp = Input(shape=(30, ))
        x = Dense(256, activation='relu')(inp)
        x = Dense(128, activation='relu')(x)
        x = Dropout(0.5)(x)
        x = Dense(128, activation='relu')(x)
        x = Dense(128, activation='relu')(x)
        x = Dropout(0.5)(x)
        x = Dense(7, activation='softmax')(x)

        nn = Model(inputs=inp, outputs=x)
        nn.compile(loss=SparseCategoricalCrossentropy(), optimizer=Adamax(learning_rate=0.001), metrics=[SparseCategoricalAccuracy()])
        nn.summary()
```

Model: "functional"

Layer (type)	Output Shape
input_layer (InputLayer)	(None, 30)
dense (Dense)	(None, 256)
dense_1 (Dense)	(None, 128)
dropout (Dropout)	(None, 128)
dense_2 (Dense)	(None, 128)
dense_3 (Dense)	(None, 128)
dropout_1 (Dropout)	(None, 128)
dense_4 (Dense)	(None, 7)

Total params: 74,759 (292.03 KB)

Trainable params: 74,759 (292.03 KB)

Non-trainable params: 0 (0.00 B)

```
In [ ]: epochs = 1000
        history = nn.fit(train_data,
                          epochs=epochs,
                          verbose=1,
                          validation_data=val_data,
                          class_weight=class_weight,
                          callbacks=[early]
                          )
```

Epoch 1/1000
48/48  **3s** 8ms/step - loss: 1.8327 - sparse_categorical_accuracy: 0.2595 - val_loss: 1.2738 - val_sparse_categorical_accuracy: 0.4511
Epoch 2/1000
48/48  **0s** 5ms/step - loss: 1.3165 - sparse_categorical_accuracy: 0.4592 - val_loss: 1.0407 - val_sparse_categorical_accuracy: 0.6617
Epoch 3/1000
48/48  **0s** 4ms/step - loss: 1.1631 - sparse_categorical_accuracy: 0.5468 - val_loss: 0.9571 - val_sparse_categorical_accuracy: 0.6692
Epoch 4/1000
48/48  **0s** 5ms/step - loss: 1.0944 - sparse_categorical_accuracy: 0.5979 - val_loss: 0.9358 - val_sparse_categorical_accuracy: 0.6654
Epoch 5/1000
48/48  **0s** 4ms/step - loss: 0.9901 - sparse_categorical_accuracy: 0.6048 - val_loss: 0.9158 - val_sparse_categorical_accuracy: 0.6617
Epoch 6/1000
48/48  **0s** 5ms/step - loss: 0.9382 - sparse_categorical_accuracy: 0.6409 - val_loss: 0.8158 - val_sparse_categorical_accuracy: 0.6992
Epoch 7/1000
48/48  **0s** 5ms/step - loss: 0.8868 - sparse_categorical_accuracy: 0.6698 - val_loss: 0.8139 - val_sparse_categorical_accuracy: 0.7293
Epoch 8/1000
48/48  **0s** 5ms/step - loss: 0.7843 - sparse_categorical_accuracy: 0.7220 - val_loss: 0.7592 - val_sparse_categorical_accuracy: 0.7556
Epoch 9/1000
48/48  **0s** 5ms/step - loss: 0.7551 - sparse_categorical_accuracy: 0.7096 - val_loss: 0.7370 - val_sparse_categorical_accuracy: 0.7406
Epoch 10/1000
48/48  **0s** 5ms/step - loss: 0.7274 - sparse_categorical_accuracy: 0.7371 - val_loss: 0.7488 - val_sparse_categorical_accuracy: 0.7519
Epoch 11/1000
48/48  **0s** 4ms/step - loss: 0.6668 - sparse_categorical_accuracy: 0.7442 - val_loss: 0.7169 - val_sparse_categorical_accuracy: 0.7744
Epoch 12/1000
48/48  **0s** 4ms/step - loss: 0.6312 - sparse_categorical_accuracy: 0.7730 - val_loss: 0.7175 - val_sparse_categorical_accuracy: 0.7895
Epoch 13/1000
48/48  **0s** 5ms/step - loss: 0.6090 - sparse_categorical_accuracy: 0.7854 - val_loss: 0.7165 - val_sparse_categorical_accuracy: 0.7782
Epoch 14/1000
48/48  **0s** 4ms/step - loss: 0.5573 - sparse_categorical_accuracy: 0.8115 - val_loss: 0.7084 - val_sparse_categorical_accuracy: 0.8045
Epoch 15/1000
48/48  **0s** 4ms/step - loss: 0.5673 - sparse_categorical_accuracy: 0.7960 - val_loss: 0.7481 - val_sparse_categorical_accuracy: 0.7895
Epoch 16/1000
48/48  **0s** 5ms/step - loss: 0.5021 - sparse_categorical_accuracy: 0.8189 - val_loss: 0.7435 - val_sparse_categorical_accuracy: 0.7632
Epoch 17/1000
48/48  **0s** 5ms/step - loss: 0.4850 - sparse_categorical_accuracy: 0.8306 - val_loss: 0.7140 - val_sparse_categorical_accuracy: 0.8083
Epoch 18/1000
48/48  **0s** 4ms/step - loss: 0.4586 - sparse_categorical_accuracy: 0.8487 - val_loss: 0.7559 - val_sparse_categorical_accuracy: 0.8045
Epoch 19/1000
48/48  **0s** 4ms/step - loss: 0.4490 - sparse_categorical_accuracy: 0.8487 - val_loss: 0.7559 - val_sparse_categorical_accuracy: 0.8045


```
ccuracy: 0.8545 - val_loss: 0.8031 - val_sparse_categorical_accuracy: 0.8083
Epoch 20/1000
48/48 ————— 0s 5ms/step - loss: 0.4681 - sparse_categorical_a
ccuracy: 0.8568 - val_loss: 0.7520 - val_sparse_categorical_accuracy: 0.8083
Epoch 21/1000
48/48 ————— 0s 5ms/step - loss: 0.4161 - sparse_categorical_a
ccuracy: 0.8657 - val_loss: 0.7829 - val_sparse_categorical_accuracy: 0.7970
Epoch 22/1000
48/48 ————— 0s 5ms/step - loss: 0.4019 - sparse_categorical_a
ccuracy: 0.8610 - val_loss: 0.8005 - val_sparse_categorical_accuracy: 0.7932
Epoch 23/1000
48/48 ————— 0s 4ms/step - loss: 0.4139 - sparse_categorical_a
ccuracy: 0.8593 - val_loss: 0.8020 - val_sparse_categorical_accuracy: 0.7932
Epoch 24/1000
48/48 ————— 0s 5ms/step - loss: 0.3586 - sparse_categorical_a
ccuracy: 0.8893 - val_loss: 0.8445 - val_sparse_categorical_accuracy: 0.7970
```

```
In [ ]: y_pred_test_nn = nn.predict(test_data, batch_size=batch_size)

roc_auc_score(y_test, y_pred_test_nn, multi_class='ovr')
```

```
10/10 ————— 0s 9ms/step
```

```
Out[ ]: 0.9612609703428328
```

Comparison of models

```
In [ ]: from sklearn.metrics import f1_score, accuracy_score, precision_score, recall_score

def get_scores(y_true, y_pred, y_pred_proba):
    return {'AUC_ROC':roc_auc_score(y_true, y_pred_proba, multi_class='ovr'),
            'F1_Score':f1_score(y_true, y_pred, average='weighted'),
            'Accuracy':accuracy_score(y_true, y_pred),
            'Precision':precision_score(y_true, y_pred, average='weighted'),
            'Recall':recall_score(y_true, y_pred, average='weighted')}

scores = []
for model in [gs_lr, gs_dt, gs_rf, xgb_clf, cbc, nn]:
    if model == xgb_clf:
        y_test_pred = np.argmax(model.predict(dtest), axis=1)
        y_test_pred_proba = model.predict(dtest)
    elif model == nn:
        y_test_pred = np.argmax(model.predict(test_data), axis=1)
        y_test_pred_proba = model.predict(test_data)
    else:
        y_test_pred = model.predict(X_test)
        y_test_pred_proba = model.predict_proba(X_test)
    scores.append(get_scores(y_test, y_test_pred, y_test_pred_proba).values())

comparison = pd.DataFrame(data=scores, index=["Logistic Regression", "Decision Tree", "Random Forest", "XGBoost", "Cubic", "Neural Network"])
comparison.style.highlight_max(color='green', axis=0).highlight_min(color='red')
```

```
10/10 ————— 0s 2ms/step
```

```
10/10 ————— 0s 2ms/step
```

Out []:

	AUC_ROC	F1_Score	Accuracy	Precision	Recall
Logistic Regression	0.889393	0.621080	0.636943	0.631468	0.636943
Decision Tree	0.921303	0.682760	0.681529	0.691372	0.681529
Random Forest	0.972110	0.838787	0.840764	0.842998	0.840764
XGBoost	0.980677	0.866181	0.866242	0.867547	0.866242
CatBoost	0.984935	0.891044	0.891720	0.891502	0.891720
Neural Network	0.961261	0.831495	0.831210	0.835182	0.831210

Confusion matrix

```
In [ ]: from sklearn.metrics import confusion_matrix

def my_cm(y_true, y_pred, title):
    cm_val = confusion_matrix(y_true, y_pred)
    cm_pgs = np.round(confusion_matrix(y_true, y_pred, normalize='true')*100


    formatted_text = (np.asarray([f"{pgs}%\n({val})" for val, pgs in zip(cm_

    sns.heatmap(cm_pgs, annot=formatted_text, fmt='', cmap='BuPu', yticklabe
    plt.title(title)
    plt.xlabel("Prediction")
    plt.ylabel("Expected")

    plt.subplots_adjust(hspace=0.5)
    return

y_test_pred_cbc = cbc.predict(X_test)

plt.figure(figsize=(10, 10))
my_cm(y_test, y_test_pred_cbc, title="Catboost")
```




Feature importance analysis with SHAP

SHAP is a method that explains how individual predictions are made by a machine learning model. SHAP deconstructs a prediction into a sum of contributions from each of the model's input variables. For each instance in the data, the contribution from each input variable towards the model's prediction will vary depending on the values of the variables for that particular instance.

A machine learning model's prediction, $f(x)$, can be represented as the sum of its computed SHAP values, plus a fixed base value, such that:

$$f(x) = \text{base value} + \text{SUM}(\text{SHAP values})$$



```
In [ ]: import shap
```

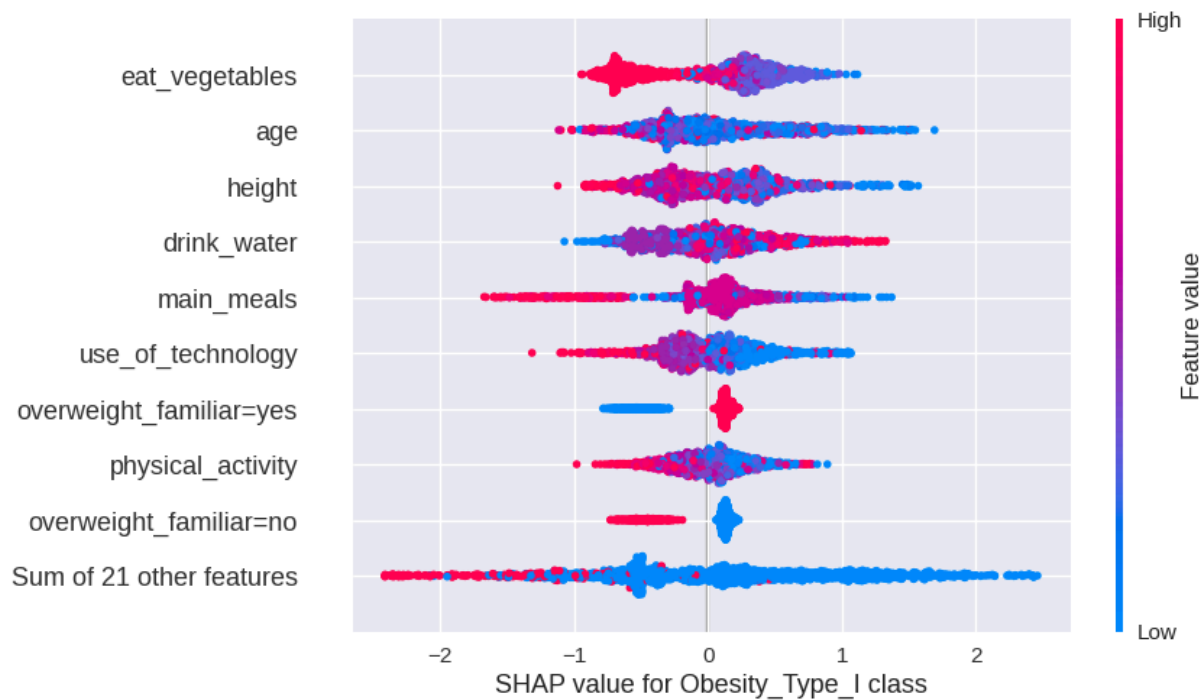
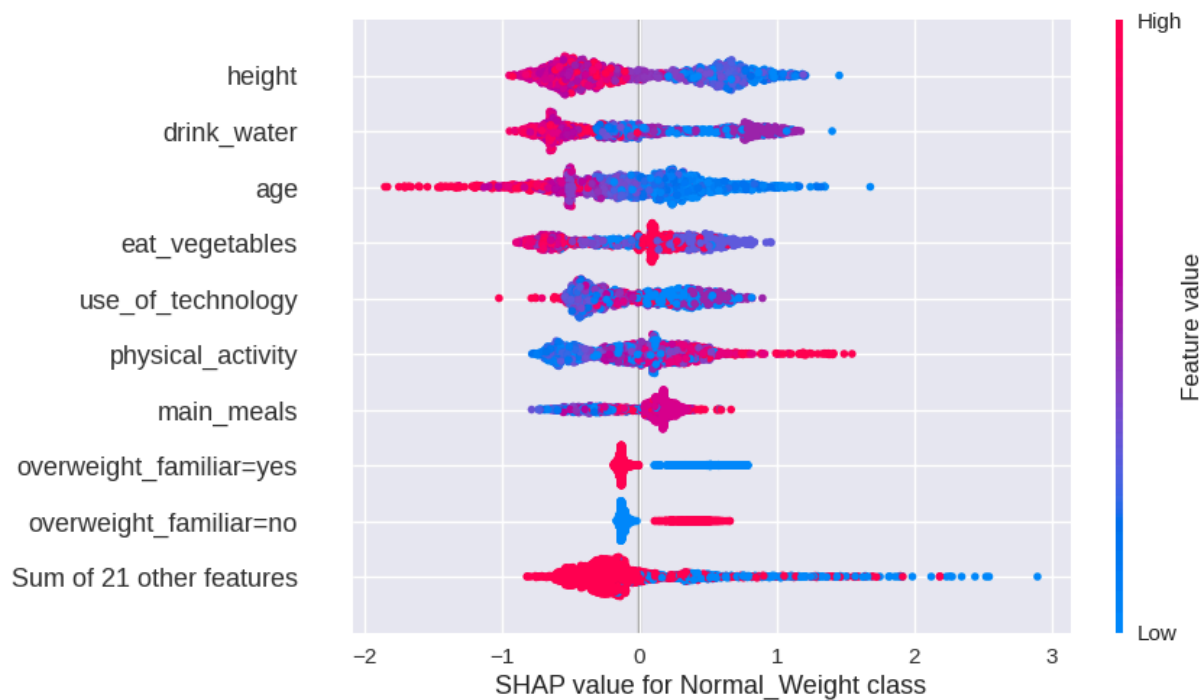
```
shap.initjs()
```




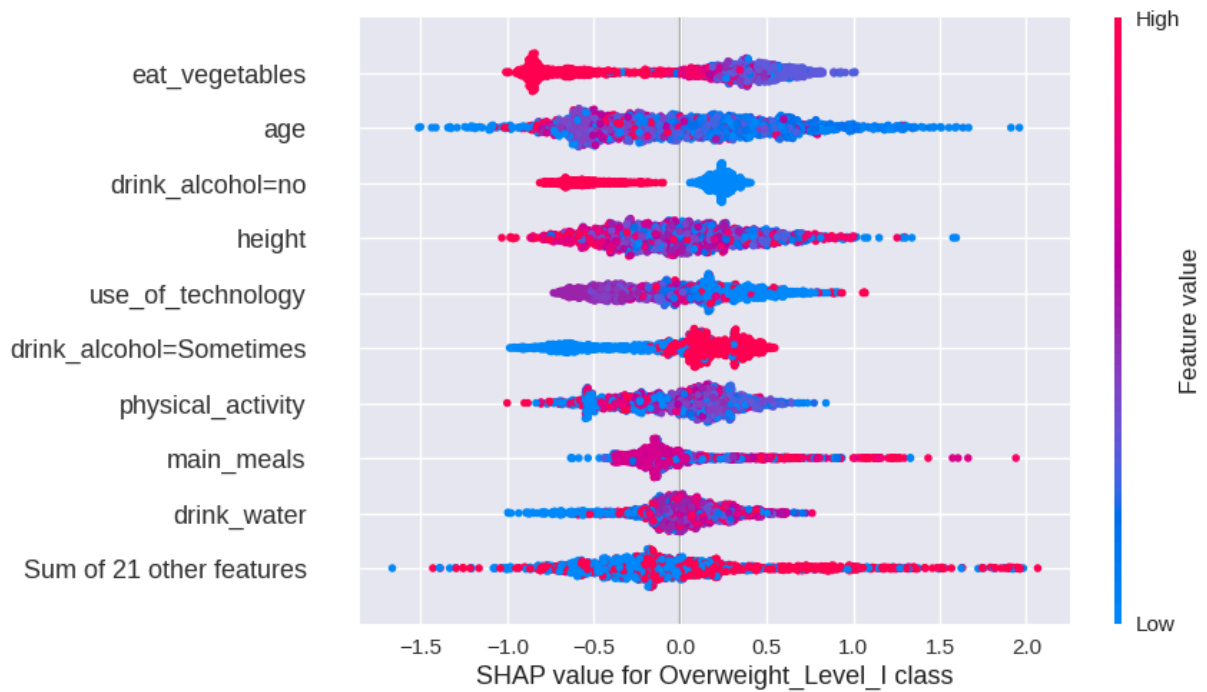
```
In [ ]: explainer = shap.TreeExplainer(cbc)
shap_values = explainer(X_train, y_train)
```

```
In [ ]: plt.style.use("seaborn-v0_8")
for i in range(6):
    shap.plots.beeswarm(shap_values[..., i], show=False)
    plt.xlabel(f"SHAP value for {le.classes_[i]} class")
    plt.show()
```









According to the diagram presented above, when a point (an instance) is blue, it indicates that its original feature value in the dataset is low, whereas when a point is red, it indicates that the original feature value in the dataset is high.

With this clarified, four situations can arise:

- Red point scattered towards negative SHAP values: For this instance, the feature indicates that the higher its value, the lower the prediction value, as it requires the SHAP value contribution to the prediction to be smaller.
- Red point scattered towards positive SHAP values: On the contrary, for this instance, the feature indicates that the higher its value, the higher the prediction value, as it requires the SHAP value contribution to the prediction to be larger.
- Blue point scattered towards negative SHAP values: For this instance, the feature indicates that the lower its value, the lower the prediction value, as it requires the SHAP value contribution to the prediction to be smaller.
- Blue point scattered towards positive SHAP values: For this instance, the feature indicates that the lower its value, the higher the prediction value, as it requires the SHAP value contribution to the prediction to be larger.

In this way, we can conclude that, for example, for the Insufficient_Weight class, the individual's age, the higher it is, the less likely it is that this individual belongs to this class. On the other hand, it can be seen that for the Normal_Weight class, it is more likely that individuals who belong to this class have high physical activity.

Saving the pipeline, the label encoder and the model with CloudPickle

An important difference between `cloudpickle` and `pickle` is that `cloudpickle` can serialize a function or class **by value**, whereas `pickle` can only serialize it **by reference**.

Serialization by reference treats functions and classes as attributes of modules, and pickles them through instructions that trigger the import of their module at load time.

Serialization by reference is thus limited in that it assumes that the module containing the function or class is available/importable in the unpickling environment. This assumption breaks when pickling constructs defined in an interactive session, a case that is automatically detected by `cloudpickle`, that pickles such constructs **by value**.

```
In [ ]: import cloudpickle

with open('obesity-levels-model.bin', 'wb') as f_out:
    cloudpickle.dump((pipe, le, cbc), f_out)
```

Converting Notebook to PDF

```
In [1]: !apt-get install texlive-xetex texlive-fonts-recommended texlive-plain-generic
!pip install py pandoc nbconvert[webpdf]
!playwright install

from google.colab import drive
from IPython.display import clear_output
drive.mount('/content/drive')
clear_output(wait=False)
```

```
In [4]: %%capture
!jupyter nbconvert --to webpdf /content/drive/MyDrive/ML_Zoomcamps_2024/Midt
```